

```
# Install required libraries
!pip install transformers # Installs the transformers library for NLP models
!pip install imageio # Installs imageio to handle image files (e.g., GIFs)
!pip install gtts # Installs Google Text-to-Speech API for text-to-speech
!pip install bert-score # Installs the BERTScore metric library for evaluation
!pip install nltk # Installs the NLTK library for natural language processing tasks
!pip install rouge-score # Installs the Rouge score package for evaluation metrics
```

```
import os
import json
import random
import requests
import imageio
from PIL import Image
from io import BytesIO
import torch
from transformers import (
    BlipProcessor,
    BlipForConditionalGeneration,
    T5Tokenizer,
    T5ForConditionalGeneration,
    AdamW
)
from torch.utils.data import Dataset, DataLoader, random_split
from tqdm import tqdm
import nltk
from nltk.translate.bleu_score import sentence_bleu, corpus_bleu
# from nltk.translate.meteor_score import meteor_score
from rouge_score import rouge_scorer
from bert_score import score as bert_score_fn
import numpy as np
```

```
# Download NLTK data
nltk.download('wordnet')
```

```
Requirement already satisfied: transformers in
/usr/local/lib/python3.10/dist-packages (4.44.2)
Requirement already satisfied: filelock in
/usr/local/lib/python3.10/dist-packages (from transformers) (3.16.1)
Requirement already satisfied: huggingface-hub<1.0,>=0.23.2 in
/usr/local/lib/python3.10/dist-packages (from transformers) (0.24.7)
Requirement already satisfied: numpy>=1.17 in
/usr/local/lib/python3.10/dist-packages (from transformers) (1.26.4)
Requirement already satisfied: packaging>=20.0 in
```

/usr/local/lib/python3.10/dist-packages (from transformers) (24.1)  
Requirement already satisfied: pyyaml<=5.1 in  
/usr/local/lib/python3.10/dist-packages (from transformers) (6.0.2)  
Requirement already satisfied: regex!=2019.12.17 in  
/usr/local/lib/python3.10/dist-packages (from transformers)  
(2024.9.11)  
Requirement already satisfied: requests in  
/usr/local/lib/python3.10/dist-packages (from transformers) (2.32.3)  
Requirement already satisfied: safetensors<=0.4.1 in  
/usr/local/lib/python3.10/dist-packages (from transformers) (0.4.5)  
Requirement already satisfied: tokenizers<0.20,>=0.19 in  
/usr/local/lib/python3.10/dist-packages (from transformers) (0.19.1)  
Requirement already satisfied: tqdm<=4.27 in  
/usr/local/lib/python3.10/dist-packages (from transformers) (4.66.6)  
Requirement already satisfied: fsspec<=2023.5.0 in  
/usr/local/lib/python3.10/dist-packages (from huggingface-  
hub<1.0,>=0.23.2->transformers) (2024.10.0)  
Requirement already satisfied: typing-extensions<=3.7.4.3 in  
/usr/local/lib/python3.10/dist-packages (from huggingface-  
hub<1.0,>=0.23.2->transformers) (4.12.2)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.10/dist-packages (from requests->transformers)  
(3.4.0)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.10/dist-packages (from requests->transformers)  
(3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.10/dist-packages (from requests->transformers)  
(2.2.3)  
Requirement already satisfied: certifi<=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests->transformers)  
(2024.8.30)  
Requirement already satisfied: imageio in  
/usr/local/lib/python3.10/dist-packages (2.36.0)  
Requirement already satisfied: numpy in  
/usr/local/lib/python3.10/dist-packages (from imageio) (1.26.4)  
Requirement already satisfied: pillow<=8.3.2 in  
/usr/local/lib/python3.10/dist-packages (from imageio) (10.4.0)  
Collecting gtts  
 Downloading gtts-2.5.3-py3-none-any.whl.metadata (4.1 kB)  
Requirement already satisfied: requests<3,>=2.27 in  
/usr/local/lib/python3.10/dist-packages (from gtts) (2.32.3)  
Requirement already satisfied: click<8.2,>=7.1 in  
/usr/local/lib/python3.10/dist-packages (from gtts) (8.1.7)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27->gtts)  
(3.4.0)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27->gtts)

(3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27->gtts)  
(2.2.3)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.27->gtts)  
(2024.8.30)  
Downloading gTTS-2.5.3-py3-none-any.whl (29 kB)  
Installing collected packages: gtts  
Successfully installed gtts-2.5.3  
Collecting bert-score  
  Downloading bert\_score-0.3.13-py3-none-any.whl.metadata (15 kB)  
Requirement already satisfied: torch>=1.0.0 in  
/usr/local/lib/python3.10/dist-packages (from bert-score)  
(2.5.0+cu121)  
Requirement already satisfied: pandas>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (2.2.2)  
Requirement already satisfied: transformers>=3.0.0 in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (4.44.2)  
Requirement already satisfied: numpy in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (1.26.4)  
Requirement already satisfied: requests in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (2.32.3)  
Requirement already satisfied: tqdm>=4.31.1 in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (4.66.6)  
Requirement already satisfied: matplotlib in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (3.8.0)  
Requirement already satisfied: packaging>=20.9 in  
/usr/local/lib/python3.10/dist-packages (from bert-score) (24.1)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.1->bert-  
score) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in  
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.1->bert-  
score) (2024.2)  
Requirement already satisfied: tzdata>=2022.7 in  
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.1->bert-  
score) (2024.2)  
Requirement already satisfied: filelock in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-  
score) (3.16.1)  
Requirement already satisfied: typing-extensions>=4.8.0 in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-  
score) (4.12.2)  
Requirement already satisfied: networkx in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-  
score) (3.4.2)  
Requirement already satisfied: jinja2 in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-

score) (3.1.4)  
Requirement already satisfied: fsspec in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-  
score) (2024.10.0)  
Requirement already satisfied: sympy==1.13.1 in  
/usr/local/lib/python3.10/dist-packages (from torch>=1.0.0->bert-  
score) (1.13.1)  
Requirement already satisfied: mpmath<1.4,>=1.1.0 in  
/usr/local/lib/python3.10/dist-packages (from sympy==1.13.1-  
>torch>=1.0.0->bert-score) (1.3.0)  
Requirement already satisfied: huggingface-hub<1.0,>=0.23.2 in  
/usr/local/lib/python3.10/dist-packages (from transformers>=3.0.0-  
>bert-score) (0.24.7)  
Requirement already satisfied: pyyaml>=5.1 in  
/usr/local/lib/python3.10/dist-packages (from transformers>=3.0.0-  
>bert-score) (6.0.2)  
Requirement already satisfied: regex!=2019.12.17 in  
/usr/local/lib/python3.10/dist-packages (from transformers>=3.0.0-  
>bert-score) (2024.9.11)  
Requirement already satisfied: safetensors>=0.4.1 in  
/usr/local/lib/python3.10/dist-packages (from transformers>=3.0.0-  
>bert-score) (0.4.5)  
Requirement already satisfied: tokenizers<0.20,>=0.19 in  
/usr/local/lib/python3.10/dist-packages (from transformers>=3.0.0-  
>bert-score) (0.19.1)  
Requirement already satisfied: contourpy>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(1.3.0)  
Requirement already satisfied: cycler>=0.10 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(4.54.1)  
Requirement already satisfied: kiwisolver>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(1.4.7)  
Requirement already satisfied: pillow>=6.2.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(10.4.0)  
Requirement already satisfied: pyparsing>=2.3.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib->bert-score)  
(3.2.0)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.10/dist-packages (from requests->bert-score)  
(3.4.0)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.10/dist-packages (from requests->bert-score)  
(3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.10/dist-packages (from requests->bert-score)  
(2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests->bert-score)  
(2024.8.30)

Requirement already satisfied: six>=1.5 in  
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2-  
>pandas>=1.0.1->bert-score) (1.16.0)

Requirement already satisfied: MarkupSafe>=2.0 in  
/usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.0.0-  
>bert-score) (3.0.2)

Downloading bert\_score-0.3.13-py3-none-any.whl (61 kB)  
61.1/61.1 kB 4.4 MB/s eta

0:00:00

Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages  
(3.8.1)

Requirement already satisfied: click in  
/usr/local/lib/python3.10/dist-packages (from nltk) (8.1.7)

Requirement already satisfied: joblib in  
/usr/local/lib/python3.10/dist-packages (from nltk) (1.4.2)

Requirement already satisfied: regex>=2021.8.3 in  
/usr/local/lib/python3.10/dist-packages (from nltk) (2024.9.11)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-  
packages (from nltk) (4.66.6)

Collecting rouge-score  
 Downloading rouge\_score-0.1.2.tar.gz (17 kB)  
 Preparing metadata (setup.py) ... Requirement already satisfied: absl-py  
in /usr/local/lib/python3.10/dist-packages (from rouge-score) (1.4.0)  
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-  
packages (from rouge-score) (3.8.1)  
Requirement already satisfied: numpy in  
/usr/local/lib/python3.10/dist-packages (from rouge-score) (1.26.4)  
Requirement already satisfied: six>=1.14.0 in  
/usr/local/lib/python3.10/dist-packages (from rouge-score) (1.16.0)  
Requirement already satisfied: click in  
/usr/local/lib/python3.10/dist-packages (from nltk->rouge-score)  
(8.1.7)  
Requirement already satisfied: joblib in  
/usr/local/lib/python3.10/dist-packages (from nltk->rouge-score)  
(1.4.2)  
Requirement already satisfied: regex>=2021.8.3 in  
/usr/local/lib/python3.10/dist-packages (from nltk->rouge-score)  
(2024.9.11)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-  
packages (from nltk->rouge-score) (4.66.6)  
Building wheels for collected packages: rouge-score  
 Building wheel for rouge-score (setup.py) ... e=rouge\_score-0.1.2-  
py3-none-any.whl size=24935

```
sha256=b0abd92c27af592126f598e236c66c197e64b11b92ac2650105910a37fbda83f
```

```
    Stored in directory:  
/root/.cache/pip/wheels/5f/dd/89/461065a73be61a532ff8599a28e9beef17985c9e9c31e541b4  
Successfully built rouge-score  
Installing collected packages: rouge-score  
Successfully installed rouge-score-0.1.2
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
```

```
True
```

```
import os  
import json  
import random  
  
# Define paths  
metadata_file = 'datafile.json' # Path to your metadata file that  
contains GIF information  
gifs_dir = 'gifs_temp'          # Directory to temporarily store GIFs  
  
# Create temporary GIFs directory if it doesn't exist  
os.makedirs(gifs_dir, exist_ok=True)  
  
# Load metadata from the metadata file  
with open(metadata_file, 'r') as f:  
    metadata = json.load(f)  
  
# Check if the dataset contains at least 100 GIFs, if not raise an  
error  
if len(metadata) < 1500:  
    raise ValueError("The metadata file contains fewer than 100  
GIFs.")  
  
# Randomly select 1500 GIFs from the metadata for processing  
selected_gifs = random.sample(metadata, 1500)  
  
# Output the number of selected GIFs  
print(f"Selected {len(selected_gifs)} random GIFs for processing.")  
  
import time  
  
# Enhanced download function with retries  
def download_gif(url, save_path, max_retries=3, backoff_factor=2):  
    for attempt in range(1, max_retries + 1):  
        try:  
            headers = {  
                "User-Agent": "Mozilla/5.0 (Windows NT 10.0; Win64;  
x64)"  
            }  
            response = requests.get(url, headers=headers, timeout=10)  
            response.raise_for_status()  
            with open(save_path, 'wb') as f:  
                f.write(response.content)  
            return save_path  
        except requests.exceptions.RequestException as e:  
            if attempt == max_retries:  
                raise e  
            time.sleep(backoff_factor ** attempt)
```

```

        response = requests.get(url, headers=headers, timeout=10)
        response.raise_for_status()
        with open(save_path, 'wb') as f:
            f.write(response.content)
        print(f"Downloaded: {save_path}")
        return True
    except requests.exceptions.HTTPError as http_err:
        print(f"HTTP error: {http_err} - Attempt {attempt}")
        if response.status_code == 404:
            break
    except requests.exceptions.RequestException as req_err:
        print(f"Request error: {req_err} - Attempt {attempt}")
    except Exception as e:
        print(f"Unexpected error: {e} - Attempt {attempt}")
    # Exponential backoff
    time.sleep(backoff_factor ** attempt)
    print(f"Failed to download {url}")
    return False

# Function to extract and preprocess frames from GIF
def extract_frames(gif_path, num_frames=10, frame_size=(256, 256)):
    try:
        gif = imageio.mimread(gif_path)
        total_frames = len(gif)

        if total_frames == 0:
            raise ValueError("No frames found in GIF.")

        interval = max(total_frames // num_frames, 1)
        selected_frames = [gif[i] for i in range(0, total_frames,
interval)][0:num_frames]

        processed_frames = []
        for frame in selected_frames:
            img = Image.fromarray(frame)
            if img.mode != 'RGB':
                img = img.convert('RGB')
            img_resized = img.resize(frame_size)
            processed_frames.append(np.array(img_resized))

        # Pad with the last frame or zeros
        while len(processed_frames) < num_frames:
            processed_frames.append(processed_frames[-1] if
processed_frames else np.zeros((frame_size[0], frame_size[1], 3),
dtype=np.uint8))

        return processed_frames

    except Exception as e:
        print(f"Failed to extract frames from {gif_path}: {e}")

```

```

        return []

# Function to delete GIF after processing
def delete_gif(gif_path):
    try:
        os.remove(gif_path)
        print(f"Deleted: {gif_path}")
    except Exception as e:
        print(f"Failed to delete {gif_path}: {e}")

import torch
import numpy as np
from tqdm import tqdm
import os

# Initialize lists to store data
target_frames = []
target_texts = []

# Process each GIF
for gif in tqdm(selected_gifs, desc="Processing GIFs"):
    gif_id = gif.get('id')
    url = gif.get('url')
    reference_description = gif.get('description')

    # Validate GIF entry
    if not gif_id or not url or not reference_description:
        print(f"Invalid GIF entry: {gif}")
        continue

    # Define local path to save the GIF temporarily
    gif_filename = f"{gif_id}.gif"
    gif_path = os.path.join(gifs_dir, gif_filename)

    # Download the GIF
    success = download_gif(url, gif_path)
    if not success:
        print(f"Skipping GIF {gif_id} due to download failure.")
        continue

    # Extract frames
    selected_frames = extract_frames(gif_path, num_frames=5)

    # Validate frame extraction
    if not selected_frames:
        print(f"No frames extracted for {gif_id}. Skipping.")
        delete_gif(gif_path)
        continue

    # Append the extracted frames and corresponding text

```



```

target_frames.append(selected_frames)
target_texts.append(reference_description)

# Delete the downloaded GIF
delete_gif(gif_path)

import torch
import numpy as np

# Assume target_frames is the list of GIFs, where each GIF is a list of frames (each frame is a numpy array)
processed_gifs = []

for gif in target_frames:
    # Convert list of frames (numpy arrays) into a single numpy array
    frames_array = np.array(gif) # Convert to a single numpy array
    # Now convert the numpy array to a PyTorch tensor
    frames_tensor = torch.tensor(frames_array) # Convert to a PyTorch tensor

    # Append the processed tensor to the list
    processed_gifs.append(frames_tensor)

# Now processed_gifs is a list of tensors, where each tensor represents a GIF

```

Saving the data

```

import torch

# Move all tensors in processed_gifs to CPU
processed_gifs_cpu = [gif.cpu() for gif in processed_gifs]

# Create a dictionary to hold both lists
data = {
    'processed_gifs': processed_gifs_cpu,
    'target_texts': target_texts
}

# Save the dictionary to a file
torch.save(data, 'data.pth')

print("Data has been saved successfully as 'data.pth'.")

```

Loading data when needed

```

import torch

# Load the saved data from the local file
data = torch.load('data.pth', map_location=torch.device('cpu')) #

```

*Change 'cpu' to 'cuda' if loading on GPU*

```
processed_gifs = data['processed_gifs']  
target_texts = data['target_texts']
```

```
print("Data has been loaded successfully from 'data.pth'.")
```

```
<ipython-input-10-1e2442dfdc7f>:4: FutureWarning: You are using  
'torch.load' with 'weights_only=False' (the current default value),  
which uses the default pickle module implicitly. It is possible to  
construct malicious pickle data which will execute arbitrary code  
during unpickling (See  
https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models  
for more details). In a future release, the default value for  
'weights_only' will be flipped to 'True'. This limits the functions  
that could be executed during unpickling. Arbitrary objects will no  
longer be allowed to be loaded via this mode unless they are  
explicitly allowlisted by the user via  
'torch.serialization.add_safe_globals'. We recommend you start setting  
'weights_only=True' for any use case where you don't have full control  
of the loaded file. Please open an issue on GitHub for any issues  
related to this experimental feature.
```

```
data = torch.load('data.pth', map_location=torch.device('cpu')) #  
Change 'cpu' to 'cuda' if loading on GPU
```

```
Data has been loaded successfully from 'data.pth'.
```

*# Initialize BLIP processor and model*

```
blip_processor = BlipProcessor.from_pretrained("Salesforce/blip-image-  
captioning-base")
```

```
blip_model =
```

```
BlipForConditionalGeneration.from_pretrained("Salesforce/blip-image-  
captioning-base").to('cuda' if torch.cuda.is_available() else 'cpu')
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/  
_token.py:89: UserWarning:
```

```
The secret 'HF_TOKEN' does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your  
settings tab (https://huggingface.co/settings/tokens), set it as  
secret in your Google Colab and restart your session.
```

```
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to  
access public models or datasets.
```

```
warnings.warn(  

```

```
{"model_id": "5005aed05b1241de9b429f2252fb9ecf", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "aa6fdb1b783f4538a02ca3028c9fe2b9", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id":"dec6d4407789424cbc5b9211dc7596fb","version_major":2,"version_minor":0}
```

```
{"model_id":"2bf67e30d8cc497b86005c0297381e7a","version_major":2,"version_minor":0}
```

```
{"model_id":"42c3b9c7261640bcb3b209e004e277a","version_major":2,"version_minor":0}
```

```
/usr/local/lib/python3.10/dist-packages/transformers/
tokenization_utils_base.py:1601: FutureWarning:
`clean_up_tokenization_spaces` was not set. It will be set to `True`
by default. This behavior will be deprecated in transformers v4.45, and
will be then set to `False` by default. For more details check this
issue: https://github.com/huggingface/transformers/issues/31884
warnings.warn(
```

```
{"model_id":"f1b2a252c43e45e1b1fdfbalc7cf1e22","version_major":2,"version_minor":0}
```

```
{"model_id":"ad664e9e0de64a95a9badaa7aaa00e43","version_major":2,"version_minor":0}
```

```
layers = ["layer.10", "layer.11", "layers.10", "layers.11"]
```

```
for name, param in blip_model.named_parameters():
```

```
    param.requires_grad = False
```

```
    for layer in layers:
```

```
        if layer in name:
```

```
            param.requires_grad = True
```

```
            break
```

```
for name, param in blip_model.named_parameters():
```

```
    status = "Trainable" if param.requires_grad else "Frozen"
```

```
    print(f"Layer: {name}, Status: {status}")
```

```
Layer: vision_model.embeddings.class_embedding, Status: Frozen
```

```
Layer: vision_model.embeddings.position_embedding, Status: Frozen
```

```
Layer: vision_model.embeddings.patch_embedding.weight, Status: Frozen
```

```
Layer: vision_model.embeddings.patch_embedding.bias, Status: Frozen
```

```
Layer: vision_model.encoder.layers.0.self_attn.qkv.weight, Status:
Frozen
```

```
Layer: vision_model.encoder.layers.0.self_attn.qkv.bias, Status:
Frozen
```

```
Layer: vision_model.encoder.layers.0.self_attn.projection.weight,
Status: Frozen
```

```
Layer: vision_model.encoder.layers.0.self_attn.projection.bias,
Status: Frozen
```

```
Layer: vision_model.encoder.layers.0.layer_norm1.weight, Status:
Frozen
```

```
Layer: vision_model.encoder.layers.0.layer_norm1.bias, Status: Frozen
```

```
Layer: vision_model.encoder.layers.0.mlp.fc1.weight, Status: Frozen
```

Layer: vision\_model.encoder.layers.0.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.0.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.0.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.0.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.0.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.1.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.1.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.2.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.2.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.3.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.3.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.self\_attn.projection.bias,

Status: Frozen  
Layer: vision\_model.encoder.layers.3.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.3.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.3.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.3.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.3.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.4.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.4.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.5.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.5.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.self\_attn.qkv.weight, Status: Frozen

Layer: vision\_model.encoder.layers.6.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.6.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.6.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.6.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.6.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.6.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.6.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.7.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.7.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.8.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.8.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.8.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.8.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.8.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.mlp.fc2.bias, Status: Frozen

Layer: vision\_model.encoder.layers.8.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.8.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.self\_attn.qkv.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.self\_attn.qkv.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.self\_attn.projection.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.self\_attn.projection.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.layer\_norm1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.layer\_norm1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.mlp.fc1.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.mlp.fc1.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.mlp.fc2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.mlp.fc2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.9.layer\_norm2.weight, Status: Frozen  
Layer: vision\_model.encoder.layers.9.layer\_norm2.bias, Status: Frozen  
Layer: vision\_model.encoder.layers.10.self\_attn.qkv.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.self\_attn.qkv.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.10.self\_attn.projection.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.self\_attn.projection.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.10.layer\_norm1.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.layer\_norm1.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.10.mlp.fc1.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.mlp.fc1.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.10.mlp.fc2.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.mlp.fc2.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.10.layer\_norm2.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.10.layer\_norm2.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.11.self\_attn.qkv.weight, Status: Trainable  
Layer: vision\_model.encoder.layers.11.self\_attn.qkv.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.11.self\_attn.projection.weight, Status: Trainable

Layer: vision\_model.encoder.layers.11.self\_attn.projection.bias,  
Status: Trainable  
Layer: vision\_model.encoder.layers.11.layer\_norm1.weight, Status:  
Trainable  
Layer: vision\_model.encoder.layers.11.layer\_norm1.bias, Status:  
Trainable  
Layer: vision\_model.encoder.layers.11.mlp.fc1.weight, Status:  
Trainable  
Layer: vision\_model.encoder.layers.11.mlp.fc1.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.11.mlp.fc2.weight, Status:  
Trainable  
Layer: vision\_model.encoder.layers.11.mlp.fc2.bias, Status: Trainable  
Layer: vision\_model.encoder.layers.11.layer\_norm2.weight, Status:  
Trainable  
Layer: vision\_model.encoder.layers.11.layer\_norm2.bias, Status:  
Trainable  
Layer: vision\_model.post\_layernorm.weight, Status: Frozen  
Layer: vision\_model.post\_layernorm.bias, Status: Frozen  
Layer: text\_decoder.bert.embeddings.word\_embeddings.weight, Status:  
Frozen  
Layer: text\_decoder.bert.embeddings.position\_embeddings.weight,  
Status: Frozen  
Layer: text\_decoder.bert.embeddings.LayerNorm.weight, Status: Frozen  
Layer: text\_decoder.bert.embeddings.LayerNorm.bias, Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.query.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.value.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.self.query.weight,



Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.0.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.intermediate.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.intermediate.dense.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.output.dense.weight, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.0.output.dense.bias, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.0.output.LayerNorm.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.0.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.self.query.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.self.value.weight,  
Status: Frozen

Layer: text\_decoder.bert.encoder.layer.1.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.1.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.intermediate.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.intermediate.dense.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.output.dense.weight, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.1.output.dense.bias, Status:

Frozen  
Layer: text\_decoder.bert.encoder.layer.1.output.LayerNorm.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.1.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.query.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.value.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.2.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.2.crossattention.output.dense.weight,  
Status: Frozen  
Layer:

```
text_decoder.bert.encoder.layer.2.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.2.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.2.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.2.intermediate.dense.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.2.intermediate.dense.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.2.output.dense.weight, Status:  
Frozen  
Layer: text_decoder.bert.encoder.layer.2.output.dense.bias, Status:  
Frozen  
Layer: text_decoder.bert.encoder.layer.2.output.LayerNorm.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.2.output.LayerNorm.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.query.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.query.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.key.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.key.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.value.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.self.value.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.3.attention.output.dense.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.3.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.3.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.3.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.3.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.3.crossattention.self.query.bias,  
Status: Frozen
```

```
Layer:
text_decoder.bert.encoder.layer.3.crossattention.self.key.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.3.crossattention.self.key.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.self.value.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.self.value.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.output.dense.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.output.dense.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.output.LayerNorm.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.3.crossattention.output.LayerNorm.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.3.intermediate.dense.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.3.intermediate.dense.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.3.output.dense.weight, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.3.output.dense.bias, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.3.output.LayerNorm.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.3.output.LayerNorm.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.query.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.query.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.key.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.key.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.value.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.4.attention.self.value.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.4.attention.output.dense.weight,
Status: Frozen
```

Layer: text\_decoder.bert.encoder.layer.4.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.4.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.4.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.4.intermediate.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.4.intermediate.dense.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.4.output.dense.weight, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.4.output.dense.bias, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.4.output.LayerNorm.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.4.output.LayerNorm.bias,  
Status: Frozen

Layer: text\_decoder.bert.encoder.layer.5.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.self.query.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.self.value.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.5.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.5.crossattention.output.LayerNorm.weight,  
Status: Frozen

```
Layer:
text_decoder.bert.encoder.layer.5.crossattention.output.LayerNorm.bias
, Status: Frozen
Layer: text_decoder.bert.encoder.layer.5.intermediate.dense.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.5.intermediate.dense.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.5.output.dense.weight, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.5.output.dense.bias, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.5.output.LayerNorm.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.5.output.LayerNorm.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.query.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.query.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.key.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.key.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.value.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.self.value.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.attention.output.dense.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.attention.output.dense.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.attention.output.LayerNorm.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.attention.output.LayerNorm.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.self.query.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.self.query.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.self.key.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.crossattention.self.key.bias,
Status: Frozen
```



```
Layer:
text_decoder.bert.encoder.layer.6.crossattention.self.value.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.self.value.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.output.dense.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.output.dense.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.output.LayerNorm.weight,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.6.crossattention.output.LayerNorm.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.intermediate.dense.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.intermediate.dense.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.output.dense.weight, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.6.output.dense.bias, Status:
Frozen
Layer: text_decoder.bert.encoder.layer.6.output.LayerNorm.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.6.output.LayerNorm.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.query.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.query.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.key.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.key.bias,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.value.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.self.value.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.7.attention.output.dense.weight,
Status: Frozen
Layer: text_decoder.bert.encoder.layer.7.attention.output.dense.bias,
Status: Frozen
Layer:
text_decoder.bert.encoder.layer.7.attention.output.LayerNorm.weight,
```

Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.7.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.7.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.7.intermediate.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.7.intermediate.dense.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.7.output.dense.weight, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.7.output.dense.bias, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.7.output.LayerNorm.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.7.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.self.query.bias,  
Status: Frozen

Layer: text\_decoder.bert.encoder.layer.8.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.self.value.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.self.value.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.8.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.intermediate.dense.weight,

Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.intermediate.dense.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.output.dense.weight, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.8.output.dense.bias, Status:  
Frozen  
Layer: text\_decoder.bert.encoder.layer.8.output.LayerNorm.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.8.output.LayerNorm.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.query.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.query.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.key.bias,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.value.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.self.value.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.attention.output.dense.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.attention.output.dense.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.attention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.attention.output.LayerNorm.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.crossattention.self.query.weight,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.crossattention.self.query.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.crossattention.self.key.weight,  
Status: Frozen  
Layer: text\_decoder.bert.encoder.layer.9.crossattention.self.key.bias,  
Status: Frozen  
Layer:  
text\_decoder.bert.encoder.layer.9.crossattention.self.value.weight,  
Status: Frozen  
Layer:

```
text_decoder.bert.encoder.layer.9.crossattention.self.value.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.9.crossattention.output.dense.weight,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.9.crossattention.output.dense.bias,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.9.crossattention.output.LayerNorm.weight,  
Status: Frozen  
Layer:  
text_decoder.bert.encoder.layer.9.crossattention.output.LayerNorm.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.9.intermediate.dense.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.9.intermediate.dense.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.9.output.dense.weight, Status:  
Frozen  
Layer: text_decoder.bert.encoder.layer.9.output.dense.bias, Status:  
Frozen  
Layer: text_decoder.bert.encoder.layer.9.output.LayerNorm.weight,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.9.output.LayerNorm.bias,  
Status: Frozen  
Layer: text_decoder.bert.encoder.layer.10.attention.self.query.weight,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.self.query.bias,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.self.key.weight,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.self.key.bias,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.self.value.weight,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.self.value.bias,  
Status: Trainable  
Layer:  
text_decoder.bert.encoder.layer.10.attention.output.dense.weight,  
Status: Trainable  
Layer: text_decoder.bert.encoder.layer.10.attention.output.dense.bias,  
Status: Trainable  
Layer:  
text_decoder.bert.encoder.layer.10.attention.output.LayerNorm.weight,  
Status: Trainable  
Layer:  
text_decoder.bert.encoder.layer.10.attention.output.LayerNorm.bias,  
Status: Trainable
```

```
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.query.weight,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.query.bias,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.key.weight,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.key.bias,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.value.weight,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.self.value.bias,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.output.dense.weight,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.output.dense.bias,
Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.output.LayerNorm.weight, Status: Trainable
Layer:
text_decoder.bert.encoder.layer.10.crossattention.output.LayerNorm.bias, Status: Trainable
Layer: text_decoder.bert.encoder.layer.10.intermediate.dense.weight,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.10.intermediate.dense.bias,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.10.output.dense.weight, Status:
Trainable
Layer: text_decoder.bert.encoder.layer.10.output.dense.bias, Status:
Trainable
Layer: text_decoder.bert.encoder.layer.10.output.LayerNorm.weight,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.10.output.LayerNorm.bias,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.attention.self.query.weight,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.attention.self.query.bias,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.attention.self.key.weight,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.attention.self.key.bias,
```

Status: Trainable  
Layer: text\_decoder.bert.encoder.layer.11.attention.self.value.weight,  
Status: Trainable  
Layer: text\_decoder.bert.encoder.layer.11.attention.self.value.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.attention.output.dense.weight,  
Status: Trainable  
Layer: text\_decoder.bert.encoder.layer.11.attention.output.dense.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.attention.output.LayerNorm.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.attention.output.LayerNorm.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.query.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.query.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.key.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.key.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.value.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.self.value.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.output.dense.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.output.dense.bias,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.output.LayerNorm.weight,  
Status: Trainable  
Layer:  
text\_decoder.bert.encoder.layer.11.crossattention.output.LayerNorm.bias,  
Status: Trainable  
Layer: text\_decoder.bert.encoder.layer.11.intermediate.dense.weight,  
Status: Trainable  
Layer: text\_decoder.bert.encoder.layer.11.intermediate.dense.bias,

```

Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.output.dense.weight, Status:
Trainable
Layer: text_decoder.bert.encoder.layer.11.output.dense.bias, Status:
Trainable
Layer: text_decoder.bert.encoder.layer.11.output.LayerNorm.weight,
Status: Trainable
Layer: text_decoder.bert.encoder.layer.11.output.LayerNorm.bias,
Status: Trainable
Layer: text_decoder.cls.predictions.bias, Status: Frozen
Layer: text_decoder.cls.predictions.transform.dense.weight, Status:
Frozen
Layer: text_decoder.cls.predictions.transform.dense.bias, Status:
Frozen
Layer: text_decoder.cls.predictions.transform.LayerNorm.weight,
Status: Frozen
Layer: text_decoder.cls.predictions.transform.LayerNorm.bias, Status:
Frozen
Layer: text_decoder.cls.predictions.decoder.weight, Status: Frozen

```

*# Initialize T5 tokenizer and model*

```

t5_tokenizer = T5Tokenizer.from_pretrained('t5-base')
t5_model = T5ForConditionalGeneration.from_pretrained('t5-
base').to('cuda' if torch.cuda.is_available() else 'cpu')

```

```

{"model_id": "61cf278f5bffa4f4f92dd2b5ffdf50d86", "version_major": 2, "vers
ion_minor": 0}

```

```

{"model_id": "ccfe5991560f41d38e6f58774ff7d6b7", "version_major": 2, "vers
ion_minor": 0}

```

```

{"model_id": "0e0539b3572b4ded9335203cc5345dd6", "version_major": 2, "vers
ion_minor": 0}

```

You are using the default legacy behaviour of the <class 'transformers.models.t5.tokenization\_t5.T5Tokenizer'>. This is expected, and simply means that the `legacy` (previous) behavior will be used so nothing changes for you. If you want to use the new behaviour, set `legacy=False`. This should only be set if you understand what it means, and thoroughly read the reason why this was added as explained in <https://github.com/huggingface/transformers/pull/24565>

```

{"model_id": "864e60a6b94d455fb1e2f933bc8a5480", "version_major": 2, "vers
ion_minor": 0}

```

```

{"model_id": "e744fb4be660406a9a9afef3ebf14261", "version_major": 2, "vers
ion_minor": 0}

```

```

for name, parameter in t5_model.named_parameters():
    if "encoder.block" in name:

```



```

index = name.split(".")[2]
if int(index) < 10:
    parameter.requires_grad = False
else:
    parameter.requires_grad = True

for name, parameter in t5_model.named_parameters():
    if "decoder.block" in name:
        index = name.split(".")[2]
        if int(index) < 9:
            parameter.requires_grad = False
        else:
            parameter.requires_grad = True

for name, parameter in t5_model.named_parameters():
    status = "Trainable" if parameter.requires_grad else "Frozen"
    print(f"Layer: {name}, Status: {status}")

Layer: shared.weight, Status: Trainable
Layer: encoder.block.0.layer.0.SelfAttention.q.weight, Status: Frozen
Layer: encoder.block.0.layer.0.SelfAttention.k.weight, Status: Frozen
Layer: encoder.block.0.layer.0.SelfAttention.v.weight, Status: Frozen
Layer: encoder.block.0.layer.0.SelfAttention.o.weight, Status: Frozen
Layer:
encoder.block.0.layer.0.SelfAttention.relative_attention_bias.weight,
Status: Frozen
Layer: encoder.block.0.layer.0.layer_norm.weight, Status: Frozen
Layer: encoder.block.0.layer.1.DenseReluDense.wi.weight, Status:
Frozen
Layer: encoder.block.0.layer.1.DenseReluDense.wo.weight, Status:
Frozen
Layer: encoder.block.0.layer.1.layer_norm.weight, Status: Frozen
Layer: encoder.block.1.layer.0.SelfAttention.q.weight, Status: Frozen
Layer: encoder.block.1.layer.0.SelfAttention.k.weight, Status: Frozen
Layer: encoder.block.1.layer.0.SelfAttention.v.weight, Status: Frozen
Layer: encoder.block.1.layer.0.SelfAttention.o.weight, Status: Frozen
Layer: encoder.block.1.layer.0.layer_norm.weight, Status: Frozen
Layer: encoder.block.1.layer.1.DenseReluDense.wi.weight, Status:
Frozen
Layer: encoder.block.1.layer.1.DenseReluDense.wo.weight, Status:
Frozen
Layer: encoder.block.1.layer.1.layer_norm.weight, Status: Frozen
Layer: encoder.block.2.layer.0.SelfAttention.q.weight, Status: Frozen
Layer: encoder.block.2.layer.0.SelfAttention.k.weight, Status: Frozen
Layer: encoder.block.2.layer.0.SelfAttention.v.weight, Status: Frozen
Layer: encoder.block.2.layer.0.SelfAttention.o.weight, Status: Frozen
Layer: encoder.block.2.layer.0.layer_norm.weight, Status: Frozen
Layer: encoder.block.2.layer.1.DenseReluDense.wi.weight, Status:
Frozen

```

[illegible]

Frozen  
Layer: encoder.block.7.layer.1.DenseReluDense.wo.weight, Status: Frozen  
Layer: encoder.block.7.layer.1.layer\_norm.weight, Status: Frozen  
Layer: encoder.block.8.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: encoder.block.8.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: encoder.block.8.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: encoder.block.8.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: encoder.block.8.layer.0.layer\_norm.weight, Status: Frozen  
Layer: encoder.block.8.layer.1.DenseReluDense.wi.weight, Status: Frozen  
Layer: encoder.block.8.layer.1.DenseReluDense.wo.weight, Status: Frozen  
Layer: encoder.block.8.layer.1.layer\_norm.weight, Status: Frozen  
Layer: encoder.block.9.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: encoder.block.9.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: encoder.block.9.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: encoder.block.9.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: encoder.block.9.layer.0.layer\_norm.weight, Status: Frozen  
Layer: encoder.block.9.layer.1.DenseReluDense.wi.weight, Status: Frozen  
Layer: encoder.block.9.layer.1.DenseReluDense.wo.weight, Status: Frozen  
Layer: encoder.block.9.layer.1.layer\_norm.weight, Status: Frozen  
Layer: encoder.block.10.layer.0.SelfAttention.q.weight, Status: Trainable  
Layer: encoder.block.10.layer.0.SelfAttention.k.weight, Status: Trainable  
Layer: encoder.block.10.layer.0.SelfAttention.v.weight, Status: Trainable  
Layer: encoder.block.10.layer.0.SelfAttention.o.weight, Status: Trainable  
Layer: encoder.block.10.layer.0.layer\_norm.weight, Status: Trainable  
Layer: encoder.block.10.layer.1.DenseReluDense.wi.weight, Status: Trainable  
Layer: encoder.block.10.layer.1.DenseReluDense.wo.weight, Status: Trainable  
Layer: encoder.block.10.layer.1.layer\_norm.weight, Status: Trainable  
Layer: encoder.block.11.layer.0.SelfAttention.q.weight, Status: Trainable  
Layer: encoder.block.11.layer.0.SelfAttention.k.weight, Status: Trainable  
Layer: encoder.block.11.layer.0.SelfAttention.v.weight, Status: Trainable  
Layer: encoder.block.11.layer.0.SelfAttention.o.weight, Status: Trainable  
Layer: encoder.block.11.layer.0.layer\_norm.weight, Status: Trainable  
Layer: encoder.block.11.layer.1.DenseReluDense.wi.weight, Status: Trainable

Layer: encoder.block.11.layer.1.DenseReluDense.wo.weight, Status: Trainable  
Layer: encoder.block.11.layer.1.layer\_norm.weight, Status: Trainable  
Layer: encoder.final\_layer\_norm.weight, Status: Trainable  
Layer: decoder.block.0.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.0.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.0.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.0.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.0.layer.0.SelfAttention.relative\_attention\_bias.weight, Status: Frozen  
Layer: decoder.block.0.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.0.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.0.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.0.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.0.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.0.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.0.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.0.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.0.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.1.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.1.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.1.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.1.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.1.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.1.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.1.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.1.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.1.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.1.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.1.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.1.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.1.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.2.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.2.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.2.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.2.layer.0.SelfAttention.o.weight, Status: Frozen

Layer: decoder.block.2.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.2.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.2.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.2.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.2.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.2.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.2.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.2.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.2.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.3.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.3.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.3.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.3.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.3.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.3.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.3.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.3.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.3.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.3.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.3.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.3.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.3.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.4.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.4.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.4.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.4.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.4.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.4.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.4.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.4.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.4.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.4.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.4.layer.2.DenseReluDense.wi.weight, Status: Frozen

Frozen  
Layer: decoder.block.4.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.4.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.5.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.5.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.5.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.5.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.5.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.5.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.5.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.5.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.5.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.5.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.5.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.5.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.5.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.6.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.6.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.6.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.6.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.6.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.6.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.6.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.6.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.6.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.6.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.6.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.6.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.6.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.7.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.7.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.7.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.7.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.7.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.7.layer.1.EncDecAttention.q.weight, Status: Frozen

Layer: decoder.block.7.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.7.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.7.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.7.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.7.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.7.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.7.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.8.layer.0.SelfAttention.q.weight, Status: Frozen  
Layer: decoder.block.8.layer.0.SelfAttention.k.weight, Status: Frozen  
Layer: decoder.block.8.layer.0.SelfAttention.v.weight, Status: Frozen  
Layer: decoder.block.8.layer.0.SelfAttention.o.weight, Status: Frozen  
Layer: decoder.block.8.layer.0.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.8.layer.1.EncDecAttention.q.weight, Status: Frozen  
Layer: decoder.block.8.layer.1.EncDecAttention.k.weight, Status: Frozen  
Layer: decoder.block.8.layer.1.EncDecAttention.v.weight, Status: Frozen  
Layer: decoder.block.8.layer.1.EncDecAttention.o.weight, Status: Frozen  
Layer: decoder.block.8.layer.1.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.8.layer.2.DenseReluDense.wi.weight, Status: Frozen  
Layer: decoder.block.8.layer.2.DenseReluDense.wo.weight, Status: Frozen  
Layer: decoder.block.8.layer.2.layer\_norm.weight, Status: Frozen  
Layer: decoder.block.9.layer.0.SelfAttention.q.weight, Status: Trainable  
Layer: decoder.block.9.layer.0.SelfAttention.k.weight, Status: Trainable  
Layer: decoder.block.9.layer.0.SelfAttention.v.weight, Status: Trainable  
Layer: decoder.block.9.layer.0.SelfAttention.o.weight, Status: Trainable  
Layer: decoder.block.9.layer.0.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.9.layer.1.EncDecAttention.q.weight, Status: Trainable  
Layer: decoder.block.9.layer.1.EncDecAttention.k.weight, Status: Trainable  
Layer: decoder.block.9.layer.1.EncDecAttention.v.weight, Status: Trainable  
Layer: decoder.block.9.layer.1.EncDecAttention.o.weight, Status: Trainable  
Layer: decoder.block.9.layer.1.layer\_norm.weight, Status: Trainable

Layer: decoder.block.9.layer.2.DenseReluDense.wi.weight, Status: Trainable  
Layer: decoder.block.9.layer.2.DenseReluDense.wo.weight, Status: Trainable  
Layer: decoder.block.9.layer.2.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.10.layer.0.SelfAttention.q.weight, Status: Trainable  
Layer: decoder.block.10.layer.0.SelfAttention.k.weight, Status: Trainable  
Layer: decoder.block.10.layer.0.SelfAttention.v.weight, Status: Trainable  
Layer: decoder.block.10.layer.0.SelfAttention.o.weight, Status: Trainable  
Layer: decoder.block.10.layer.0.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.10.layer.1.EncDecAttention.q.weight, Status: Trainable  
Layer: decoder.block.10.layer.1.EncDecAttention.k.weight, Status: Trainable  
Layer: decoder.block.10.layer.1.EncDecAttention.v.weight, Status: Trainable  
Layer: decoder.block.10.layer.1.EncDecAttention.o.weight, Status: Trainable  
Layer: decoder.block.10.layer.1.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.10.layer.2.DenseReluDense.wi.weight, Status: Trainable  
Layer: decoder.block.10.layer.2.DenseReluDense.wo.weight, Status: Trainable  
Layer: decoder.block.10.layer.2.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.11.layer.0.SelfAttention.q.weight, Status: Trainable  
Layer: decoder.block.11.layer.0.SelfAttention.k.weight, Status: Trainable  
Layer: decoder.block.11.layer.0.SelfAttention.v.weight, Status: Trainable  
Layer: decoder.block.11.layer.0.SelfAttention.o.weight, Status: Trainable  
Layer: decoder.block.11.layer.0.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.11.layer.1.EncDecAttention.q.weight, Status: Trainable  
Layer: decoder.block.11.layer.1.EncDecAttention.k.weight, Status: Trainable  
Layer: decoder.block.11.layer.1.EncDecAttention.v.weight, Status: Trainable  
Layer: decoder.block.11.layer.1.EncDecAttention.o.weight, Status: Trainable  
Layer: decoder.block.11.layer.1.layer\_norm.weight, Status: Trainable  
Layer: decoder.block.11.layer.2.DenseReluDense.wi.weight, Status: Trainable  
Layer: decoder.block.11.layer.2.DenseReluDense.wo.weight, Status: Trainable



Layer: decoder.block.11.layer.2.layer\_norm.weight, Status: Trainable  
Layer: decoder.final\_layer\_norm.weight, Status: Trainable

```
import torch
import numpy as np
from torch.utils.data import Dataset

class GifAggregationDataset(Dataset):
    def __init__(self, frames_list, targets, tokenizer,
max_target_length=150):
        """
        frames_list: A list of lists containing frames (as tensors or
numpy arrays) for each GIF.
        targets: The ground truth captions for each GIF.
        tokenizer: Tokenizer for T5 to encode the target captions.
        max_target_length: Maximum length for the target text
sequences.
        """
        self.frames_list = frames_list # List of frame sequences
(e.g., list of tensors for each GIF)
        self.targets = targets # Ground truth captions for each GIF
        self.tokenizer = tokenizer
        self.max_target_length = max_target_length

    def __len__(self):
        return len(self.frames_list)

    def __getitem__(self, idx):
        # Get frames and target for the given index
        frames = self.frames_list[idx] # List or tensor of frames for
a single GIF
        target_text = self.targets[idx]

        # Ensure frames are in the correct shape (num_frames, height,
width, channels)
        if isinstance(frames, list):
            # Convert list of numpy arrays to a single numpy array, if
necessary
            frames = np.array(frames)

        # If frames are numpy arrays, convert to torch tensor and make
sure dimensions are correct
        if isinstance(frames, np.ndarray):
            frames = torch.tensor(frames, dtype=torch.float32) #
Convert to float32 tensor for PyTorch

        # Reshape or permute frames to ensure the shape is
(num_frames, height, width, channels)
        if frames.ndim == 4 and frames.shape[-1] == 3:
            # Shape is (num_frames, height, width, channels), which is
```

```

correct for PIL conversion
    pass
    elif frames.ndim == 4 and frames.shape[1] == 3:
        # If frames have shape (num_frames, channels, height,
        width), change to (num_frames, height, width, channels)
        frames = frames.permute(0, 2, 3, 1) # (N, C, H, W) -> (N,
H, W, C)

    # Tokenize the target caption (ground-truth description)
    target_encoding = self.tokenizer(
        target_text,
        padding='max_length',
        truncation=True,
        max_length=self.max_target_length,
        return_tensors="pt"
    )

    # Replace padding token ID's of the labels by -100 to ignore
    them in the loss
    labels = target_encoding.input_ids.squeeze()
    labels[labels == self.tokenizer.pad_token_id] = -100

    return {
        'frames': frames, # Provide frames directly in tensor
        format (num_frames, height, width, channels)
        'labels': labels, # Provide tokenized ground truth
        captions for T5
        'reference_description': target_text
    }

import torch
import torch.nn as nn
import torch.optim as optim
from transformers import T5Tokenizer
from torch.utils.data import DataLoader, random_split

# Setup the models in training mode
blip_model.train() # BLIP is now being trained, so it's in training
mode
t5_model.train() # T5 is also being trained

# Define optimizers for BLIP and T5
blip_params = filter(lambda p: p.requires_grad,
blip_model.parameters())
t5_params = filter(lambda p: p.requires_grad, t5_model.parameters())

blip_optimizer = optim.AdamW(blip_params, lr=0.00005) # Lower
learning rate since we're training both
t5_optimizer = optim.AdamW(t5_params, lr=0.00005)

```

```

# Define the loss function for T5
criterion = nn.CrossEntropyLoss(ignore_index=-100)

# Prepare DataLoader (frames and target descriptions)

aggregation_dataset = GifAggregationDataset(
    frames_list=processed_gifs,  # This should be the frames for each
    GIF (as a list of tensors)
    targets=target_texts,        # The ground truth aggregated captions
    tokenizer=t5_tokenizer,      # Tokenizer for processing the target
    texts
    max_target_length=150        # Max length for target captions
)

# Split dataset into training and validation sets (80-20 split)
train_size = int(0.8 * len(aggregation_dataset))
val_size = len(aggregation_dataset) - train_size
train_dataset, val_dataset = random_split(aggregation_dataset,
[train_size, val_size])

# Create DataLoaders
train_dataloader = DataLoader(
    train_dataset,
    batch_size=16,  # Define the batch size that fits your GPU memory
    shuffle=True,
    num_workers=0,
    pin_memory=True
)

val_dataloader = DataLoader(
    val_dataset,
    batch_size=16,  # Define the batch size for validation
    shuffle=False,
    num_workers=0,
    pin_memory=True
)

```

Function to evaluate the model in the training phase

```

import torch
from tqdm import tqdm
from rouge_score import rouge_scorer
from nltk.translate.bleu_score import corpus_bleu
from nltk.translate.meteor_score import meteor_score
import bert_score
from PIL import Image

def evaluate_model_training(blip_model, t5_model, blip_processor,
tokenizer, dataloader, device='cuda' if torch.cuda.is_available() else
'cpu'):

```

```

blip_model.eval() # Set BLIP to evaluation mode
t5_model.eval()  # Set T5 to evaluation mode

# Metrics containers
all_references = []
all_candidates = []
rougeL_scores = []

# ROUGE scorer initialization
scorer = rouge_scorer.RougeScorer(['rougeL'], use_stemmer=True)

with torch.no_grad():
    for batch in tqdm(dataloader, desc="Evaluating"):
        # Step 1: Prepare frames for BLIP
        frames = batch['frames'] # Assuming shape [batch_size,
num_frames, C, H, W] or [batch_size, num_frames, H, W, C]
        reference_descriptions = batch['reference_description'] #
List of reference captions

        batch_size, num_frames = frames.shape[0], frames.shape[1]

        frame_captions = []

        for i in range(batch_size):
            single_captions = []
            for j in range(num_frames):
                frame = frames[i, j] # Shape depends on data
format

                # Determine frame shape and permute if necessary
                if frame.ndim == 3:
                    # Possible shapes:
                    # [C, H, W] or [H, W, C]
                    if frame.shape[0] == 3:
                        # [C, H, W] -> [H, W, C]
                        frame = frame.permute(1, 2, 0)
                    # Else assume [H, W, C], no change needed
                else:
                    raise ValueError(f"Unexpected frame shape:
{frame.shape}")

                # Convert to CPU and numpy
                frame_np = frame.cpu().numpy().astype('uint8') #
Ensure dtype is uint8

                # Convert to PIL Image
                try:
                    img = Image.fromarray(frame_np).convert('RGB')

```

```

        except Exception as e:
            print(f"Error converting frame to image: {e}")
            continue # Skip this frame

        # Generate caption using BLIP
        inputs = blip_processor(img,
return_tensors="pt").to(device)
        outputs = blip_model.generate(**inputs,
max_length=150) # Limit caption length
        caption = blip_processor.decode(outputs[0],
skip_special_tokens=True)
        single_captions.append(caption)

        # Concatenate captions for the current example
        concatenated_caption = " ".join(single_captions)
        frame_captions.append(concatenated_caption)

        # Step 2: Tokenize the concatenated captions as a batch
        max_length = 128 # Adjust based on GPU capacity and
desired sequence length
        t5_inputs = tokenizer(
            frame_captions, # List of concatenated captions, one
per example in the batch
            max_length=max_length,
            padding='max_length',
            truncation=True,
            return_tensors="pt"
        ).to(device)

        # Step 3: Generate predictions using T5
        outputs = t5_model.generate(
            input_ids=t5_inputs.input_ids,
            attention_mask=t5_inputs.attention_mask,
            max_length=150,
            num_beams=4,
            early_stopping=True
        )

        # Step 4: Decode predictions
        predictions = tokenizer.batch_decode(outputs,
skip_special_tokens=True)

        # Step 5: Evaluate the predictions
        for ref, pred in zip(reference_descriptions, predictions):
            ref_tokens = ref.split()
            pred_tokens = pred.split()

            # Append references and predictions for BLEU
            all_references.append([ref_tokens])
            all_candidates.append(pred_tokens)

```

```

        # ROUGE scores calculation
        scores = scorer.score(ref, pred)
        rougeL_scores.append(scores['rougeL'].fmeasure)

    # Calculate average ROUGE scores
    avg_rougeL = sum(rougeL_scores) / len(rougeL_scores) if
rougeL_scores else 0

    # Compile metrics
    metrics = {
        'ROUGE-L': avg_rougeL,
    }

    return metrics

```

Function to evaluate the models when they finished training phase

```

import torch
from tqdm import tqdm
from rouge_score import rouge_scorer
from nltk.translate.bleu_score import corpus_bleu
from nltk.translate.meteor_score import meteor_score
import bert_score
from PIL import Image

def evaluate_model(blip_model, t5_model, blip_processor, tokenizer,
dataloader, device='cuda' if torch.cuda.is_available() else 'cpu'):
    blip_model.eval() # Set BLIP to evaluation mode
    t5_model.eval() # Set T5 to evaluation mode

    # Metrics containers
    all_references = []
    all_candidates = []
    rouge1_scores = []
    rouge2_scores = []
    rougeL_scores = []
    # meteor_scores_list = []

    # ROUGE scorer initialization
    scorer = rouge_scorer.RougeScorer(['rouge1', 'rouge2', 'rougeL'],
use_stemmer=True)

    # BERTScore
    bert_score_fn = bert_score.score

    with torch.no_grad():
        for batch in tqdm(dataloader, desc="Evaluating"):
            # Step 1: Prepare frames for BLIP

```

```

frames = batch['frames'] # Assuming shape [batch_size,
num_frames, C, H, W] or [batch_size, num_frames, H, W, C]
reference_descriptions = batch['reference_description'] #
List of reference captions

batch_size, num_frames = frames.shape[0], frames.shape[1]

frame_captions = []

for i in range(batch_size):
    single_captions = []
    for j in range(num_frames):
        frame = frames[i, j] # Shape depends on data
format

        # Determine frame shape and permute if necessary
        if frame.ndim == 3:
            # Possible shapes:
            # [C, H, W] or [H, W, C]
            if frame.shape[0] == 3:
                # [C, H, W] -> [H, W, C]
                frame = frame.permute(1, 2, 0)
            # Else assume [H, W, C], no change needed
        else:
            raise ValueError(f"Unexpected frame shape:
{frame.shape}")

        # Convert to CPU and numpy
        frame_np = frame.cpu().numpy().astype('uint8') #
Ensure dtype is uint8

        # Convert to PIL Image
        try:
            img = Image.fromarray(frame_np).convert('RGB')
        except Exception as e:
            print(f"Error converting frame to image: {e}")
            continue # Skip this frame

        # Generate caption using BLIP
        inputs = blip_processor(img,
return_tensors="pt").to(device)
        outputs = blip_model.generate(**inputs,
max_length=50) # Limit caption length
        caption = blip_processor.decode(outputs[0],
skip_special_tokens=True)
        single_captions.append(caption)

        # Concatenate captions for the current example
        concatenated_caption = " ".join(single_captions)
        frame_captions.append(concatenated_caption)

```

```

        # Step 2: Tokenize the concatenated captions as a batch
        max_length = 128 # Adjust based on GPU capacity and
desired sequence length
        t5_inputs = tokenizer(
            frame_captions, # List of concatenated captions, one
per example in the batch
            max_length=max_length,
            padding='max_length',
            truncation=True,
            return_tensors="pt"
        ).to(device)

        # Step 3: Generate predictions using T5
        outputs = t5_model.generate(
            input_ids=t5_inputs.input_ids,
            attention_mask=t5_inputs.attention_mask,
            max_length=150,
            num_beams=4,
            early_stopping=True
        )

        # Step 4: Decode predictions
        predictions = tokenizer.batch_decode(outputs,
skip_special_tokens=True)

        # Step 5: Evaluate the predictions
        for ref, pred in zip(reference_descriptions, predictions):
            ref_tokens = ref.split()
            pred_tokens = pred.split()

            # Append references and predictions for BLEU
            all_references.append([ref_tokens])
            all_candidates.append(pred_tokens)

            # ROUGE scores calculation
            scores = scorer.score(ref, pred)
            rouge1_scores.append(scores['rouge1'].fmeasure)
            rouge2_scores.append(scores['rouge2'].fmeasure)
            rougeL_scores.append(scores['rougeL'].fmeasure)

            # METEOR score calculation
            # meteor = meteor_score([ref_tokens], pred_tokens)
            # meteor_scores_list.append(meteor)

        # Calculate BLEU score
        bleu = corpus_bleu(all_references, all_candidates)

        # Calculate average ROUGE scores
        avg_rouge1 = sum(rouge1_scores) / len(rouge1_scores) if

```



```

rouge1_scores else 0
    avg_rouge2 = sum(rouge2_scores) / len(rouge2_scores) if
rouge2_scores else 0
    avg_rougeL = sum(rougeL_scores) / len(rougeL_scores) if
rougeL_scores else 0

    # Calculate average METEOR score
    # avg_meteor = sum(meteor_scores_list) / len(meteor_scores_list)
    if meteor_scores_list else 0

    # Calculate BERT Score
    P, R, F1 = bert_score_fn(
        [' '.join(cand) for cand in all_candidates],
        [' '.join(ref[0]) for ref in all_references],
        lang='en',
        verbose=True)
    avg_bert_f1 = F1.mean().item()

    # Compile metrics
    metrics = {
        'BLEU': bleu,
        'ROUGE-1': avg_rouge1,
        'ROUGE-2': avg_rouge2,
        'ROUGE-L': avg_rougeL,
        # 'METEOR': avg_meteor,
        'BERT_F1': avg_bert_f1
    }

    return metrics

import torch
from torch.optim import AdamW
from PIL import Image
from transformers import T5Tokenizer
import torch.nn.functional as F

# Set the models to training mode
t5_model.train()
blip_model.train()

# Filter parameters that require gradients for BLIP and T5 models
blip_params = filter(lambda p: p.requires_grad,
blip_model.parameters())
t5_params = filter(lambda p: p.requires_grad, t5_model.parameters())

# Define the optimizer using AdamW with appropriate learning rates
blip_optimizer = AdamW(blip_params, lr=1e-4)
t5_optimizer = AdamW(t5_params, lr=1e-4)

# Training Loop

```

```

epochs = 7
device = 'cuda' if torch.cuda.is_available() else 'cpu'
blip_model.to(device)
t5_model.to(device)

# Lists to store metrics over epochs
training_losses = []
#bleu_scores = []
#rouge1_scores = []
#rouge2_scores = []
rougeL_scores = []
# meteor_scores = []
#bert_f1_scores = []

# Training both BLIP and T5 using the DataLoader
for epoch in range(epochs):
    t5_model.train()
    blip_model.train()

    # Initialize average loss for tracking epoch training loss
    epoch_loss = 0.0
    num_batches = 0

    for batch in train_dataloader:
        # Step 1: Prepare data for BLIP
        frames = batch['frames'] # Get frames for the current batch
        # Assuming frames is a list of shape [batch_size, num_frames,
height, width, channels]
        batch_size = frames.size(0)
        frames = frames.to(device) # Move frames to GPU if available

        # Zero gradients for both optimizers
        blip_optimizer.zero_grad()
        t5_optimizer.zero_grad()

        # Step 2: Generate frame-level captions with BLIP
        frame_captions = []
        for i in range(batch_size):
            single_captions = []
            for single_frame in frames[i]:
                # Convert frame to a PIL image and generate caption
using BLIP model
                img =
Image.fromarray(single_frame.cpu().numpy().astype('uint8'), 'RGB')
                inputs = blip_processor(img,
return_tensors="pt").to(device)

                # Set max_length explicitly for BLIP generation to
avoid inconsistent lengths
                outputs = blip_model.generate(**inputs,

```

```

max_length=200) # Limit caption length to 150 tokens
                caption = blip_processor.decode(outputs[0],
skip_special_tokens=True)
                single_captions.append(caption)
                # Concatenate captions for the current example
                concatenated_caption = " ".join(single_captions)
                frame_captions.append(concatenated_caption)

        # Step 3: Tokenize the concatenated captions as a batch
        max_length = 200 # Adjust as needed based on your GPU
capacity
        t5_inputs = t5_tokenizer(
            frame_captions, # List of concatenated captions, one per
example in the batch
            max_length=max_length,
            padding='max_length',
            truncation=True,
            return_tensors="pt"
        ).to(device)

        # Prepare labels
        labels = batch['labels'].to(device)
        labels = labels[:, :t5_inputs.input_ids.shape[-1]] # Truncate
labels if necessary

        # Pad labels to match input length using -100 to ignore during
loss calculation
        labels = F.pad(
            labels,
            (0, max(0, t5_inputs.input_ids.shape[-1] - labels.shape[-
1])),
            value=-100
        )

        # Ensure the attention mask is consistent
        t5_attention_mask = t5_inputs.attention_mask

        # Step 5: Forward pass through T5 model
        outputs = t5_model(
            input_ids=t5_inputs.input_ids,
            attention_mask=t5_attention_mask,
            labels=labels
        )
        loss = outputs.loss

        # Step 6: Backpropagation and optimizer step
        loss.backward() # Compute gradients for both models (BLIP and
T5)

        blip_optimizer.step()
        t5_optimizer.step()

```

```

    # Accumulate the loss for averaging
    epoch_loss += loss.item()
    num_batches += 1

    # Step 7: Print the training loss for the current iteration
    print(f"Epoch [{epoch + 1}/{epochs}], Batch Loss:
{loss.item()}")

    # Calculate the average training loss for the epoch
    average_epoch_loss = epoch_loss / num_batches
    training_losses.append(average_epoch_loss)
    print(f"Epoch [{epoch + 1}/{epochs}], Average Training Loss:
{average_epoch_loss}")

    # Evaluate after each epoch
    metrics = evaluate_model_training(blip_model, t5_model,
    blip_processor, t5_tokenizer, val_dataloader, device)
    #bleu_scores.append(metrics['BLEU'])
    #rouge1_scores.append(metrics['ROUGE-1'])
    #rouge2_scores.append(metrics['ROUGE-2'])
    rougeL_scores.append(metrics['ROUGE-L'])
    # meteor_scores.append(metrics['METEOR'])
    # bert_f1_scores.append(metrics['BERT_F1'])

    # Print evaluation metrics for each epoch
    print(f"Validation Metrics after Epoch {epoch+1}: {metrics}")

    # Clear GPU cache after each epoch to free up memory
    torch.cuda.empty_cache()

```

```

Epoch [1/7], Batch Loss: 2.792562484741211
Epoch [1/7], Batch Loss: 3.05653715133667
Epoch [1/7], Batch Loss: 2.8849689960479736
Epoch [1/7], Batch Loss: 2.9645094871520996
Epoch [1/7], Batch Loss: 2.6882576942443848
Epoch [1/7], Batch Loss: 2.5022499561309814
Epoch [1/7], Batch Loss: 2.851562261581421
Epoch [1/7], Batch Loss: 2.91620135307312
Epoch [1/7], Batch Loss: 3.1907904148101807
Epoch [1/7], Batch Loss: 2.9111499786376953
Epoch [1/7], Batch Loss: 3.0466115474700928
Epoch [1/7], Batch Loss: 2.8951869010925293
Epoch [1/7], Batch Loss: 2.815492868423462
Epoch [1/7], Batch Loss: 2.6365480422973633
Epoch [1/7], Batch Loss: 2.799355983734131
Epoch [1/7], Batch Loss: 2.536698579788208
Epoch [1/7], Batch Loss: 2.613626003265381
Epoch [1/7], Batch Loss: 2.661198377609253
Epoch [1/7], Batch Loss: 2.7486190795898438

```

Epoch [1/7], Batch Loss: 2.721141815185547  
Epoch [1/7], Batch Loss: 2.695596933364868  
Epoch [1/7], Batch Loss: 2.6508281230926514  
Epoch [1/7], Batch Loss: 2.557412624359131  
Epoch [1/7], Batch Loss: 3.0691192150115967  
Epoch [1/7], Batch Loss: 2.6178832054138184  
Epoch [1/7], Batch Loss: 2.716886520385742  
Epoch [1/7], Batch Loss: 2.892625331878662  
Epoch [1/7], Batch Loss: 2.5394060611724854  
Epoch [1/7], Batch Loss: 2.2684273719787598  
Epoch [1/7], Batch Loss: 2.6580679416656494  
Epoch [1/7], Batch Loss: 2.595590591430664  
Epoch [1/7], Batch Loss: 2.724113941192627  
Epoch [1/7], Batch Loss: 2.3373875617980957  
Epoch [1/7], Batch Loss: 2.622941017150879  
Epoch [1/7], Batch Loss: 2.670625925064087  
Epoch [1/7], Batch Loss: 2.3872199058532715  
Epoch [1/7], Batch Loss: 2.351562023162842  
Epoch [1/7], Batch Loss: 2.644439935684204  
Epoch [1/7], Batch Loss: 2.3942604064941406  
Epoch [1/7], Batch Loss: 2.319535493850708  
Epoch [1/7], Batch Loss: 2.723679780960083  
Epoch [1/7], Batch Loss: 2.794512987136841  
Epoch [1/7], Batch Loss: 2.486840009689331  
Epoch [1/7], Batch Loss: 2.5631191730499268  
Epoch [1/7], Batch Loss: 2.565117597579956  
Epoch [1/7], Batch Loss: 2.6146843433380127  
Epoch [1/7], Batch Loss: 2.479147434234619  
Epoch [1/7], Batch Loss: 2.392965078353882  
Epoch [1/7], Batch Loss: 2.6993730068206787  
Epoch [1/7], Batch Loss: 2.6262025833129883  
Epoch [1/7], Batch Loss: 2.429753303527832  
Epoch [1/7], Batch Loss: 2.519545316696167  
Epoch [1/7], Batch Loss: 2.65714168548584  
Epoch [1/7], Batch Loss: 2.45499587059021  
Epoch [1/7], Batch Loss: 2.672262191772461  
Epoch [1/7], Batch Loss: 2.419902801513672  
Epoch [1/7], Batch Loss: 2.4081788063049316  
Epoch [1/7], Batch Loss: 2.5227558612823486  
Epoch [1/7], Batch Loss: 2.254258632659912  
Epoch [1/7], Batch Loss: 2.55667781829834  
Epoch [1/7], Batch Loss: 2.4328277111053467  
Epoch [1/7], Batch Loss: 2.791801691055298  
Epoch [1/7], Batch Loss: 2.7886228561401367  
Epoch [1/7], Batch Loss: 2.450453281402588  
Epoch [1/7], Batch Loss: 2.238307476043701  
Epoch [1/7], Batch Loss: 2.8428313732147217  
Epoch [1/7], Batch Loss: 2.7892682552337646  
Epoch [1/7], Batch Loss: 2.473942279815674

Epoch [1/7], Batch Loss: 2.133924722671509  
Epoch [1/7], Batch Loss: 2.0994653701782227  
Epoch [1/7], Batch Loss: 2.2195956707000732  
Epoch [1/7], Batch Loss: 2.6091225147247314  
Epoch [1/7], Batch Loss: 2.849940299987793  
Epoch [1/7], Batch Loss: 2.6282076835632324  
Epoch [1/7], Batch Loss: 2.5078916549682617  
Epoch [1/7], Average Training Loss: 2.621926883061727

Evaluating: 100%|██████████| 19/19 [05:37<00:00, 17.78s/it]

Validation Metrics after Epoch 1: {'ROUGE-L': 0.3243333918983624}

Epoch [2/7], Batch Loss: 2.473907947540283  
Epoch [2/7], Batch Loss: 2.302682638168335  
Epoch [2/7], Batch Loss: 2.329249858856201  
Epoch [2/7], Batch Loss: 2.5172479152679443  
Epoch [2/7], Batch Loss: 2.761099100112915  
Epoch [2/7], Batch Loss: 2.3488035202026367  
Epoch [2/7], Batch Loss: 2.4597198963165283  
Epoch [2/7], Batch Loss: 2.6647422313690186  
Epoch [2/7], Batch Loss: 2.5523948669433594  
Epoch [2/7], Batch Loss: 2.1182148456573486  
Epoch [2/7], Batch Loss: 2.4329872131347656  
Epoch [2/7], Batch Loss: 2.316796064376831  
Epoch [2/7], Batch Loss: 2.259514808654785  
Epoch [2/7], Batch Loss: 2.6163504123687744  
Epoch [2/7], Batch Loss: 2.3733248710632324  
Epoch [2/7], Batch Loss: 2.4834394454956055  
Epoch [2/7], Batch Loss: 2.3729352951049805  
Epoch [2/7], Batch Loss: 2.58955717086792  
Epoch [2/7], Batch Loss: 2.56703782081604  
Epoch [2/7], Batch Loss: 2.5039615631103516  
Epoch [2/7], Batch Loss: 2.472188949584961  
Epoch [2/7], Batch Loss: 2.455782413482666  
Epoch [2/7], Batch Loss: 2.459120273590088  
Epoch [2/7], Batch Loss: 2.235743284225464  
Epoch [2/7], Batch Loss: 2.5564215183258057  
Epoch [2/7], Batch Loss: 2.3535382747650146  
Epoch [2/7], Batch Loss: 2.568927526473999  
Epoch [2/7], Batch Loss: 2.387378454208374  
Epoch [2/7], Batch Loss: 2.713947057723999  
Epoch [2/7], Batch Loss: 2.4272289276123047  
Epoch [2/7], Batch Loss: 2.3013620376586914  
Epoch [2/7], Batch Loss: 2.389380931854248  
Epoch [2/7], Batch Loss: 2.330120086669922  
Epoch [2/7], Batch Loss: 2.503119707107544  
Epoch [2/7], Batch Loss: 2.4521543979644775  
Epoch [2/7], Batch Loss: 2.468416213989258  
Epoch [2/7], Batch Loss: 2.0798542499542236  
Epoch [2/7], Batch Loss: 2.617144823074341

Epoch [2/7], Batch Loss: 2.2453935146331787  
Epoch [2/7], Batch Loss: 2.52299165725708  
Epoch [2/7], Batch Loss: 2.45682954788208  
Epoch [2/7], Batch Loss: 2.378944158554077  
Epoch [2/7], Batch Loss: 2.700363874435425  
Epoch [2/7], Batch Loss: 2.080040454864502  
Epoch [2/7], Batch Loss: 2.543712615966797  
Epoch [2/7], Batch Loss: 2.2072529792785645  
Epoch [2/7], Batch Loss: 2.27642822265625  
Epoch [2/7], Batch Loss: 2.521725654602051  
Epoch [2/7], Batch Loss: 2.4891200065612793  
Epoch [2/7], Batch Loss: 2.2736752033233643  
Epoch [2/7], Batch Loss: 2.6230950355529785  
Epoch [2/7], Batch Loss: 2.5693767070770264  
Epoch [2/7], Batch Loss: 2.3844621181488037  
Epoch [2/7], Batch Loss: 2.45255970954895  
Epoch [2/7], Batch Loss: 2.205230712890625  
Epoch [2/7], Batch Loss: 2.4437568187713623  
Epoch [2/7], Batch Loss: 2.4140477180480957  
Epoch [2/7], Batch Loss: 2.439300060272217  
Epoch [2/7], Batch Loss: 2.313997268676758  
Epoch [2/7], Batch Loss: 2.554961919784546  
Epoch [2/7], Batch Loss: 2.4694268703460693  
Epoch [2/7], Batch Loss: 2.4955074787139893  
Epoch [2/7], Batch Loss: 2.322580099105835  
Epoch [2/7], Batch Loss: 2.6096081733703613  
Epoch [2/7], Batch Loss: 2.311009645462036  
Epoch [2/7], Batch Loss: 2.4948084354400635  
Epoch [2/7], Batch Loss: 2.212188243865967  
Epoch [2/7], Batch Loss: 2.2408926486968994  
Epoch [2/7], Batch Loss: 2.22233510017395  
Epoch [2/7], Batch Loss: 2.5274388790130615  
Epoch [2/7], Batch Loss: 2.517465591430664  
Epoch [2/7], Batch Loss: 2.412813186645508  
Epoch [2/7], Batch Loss: 2.2280445098876953  
Epoch [2/7], Batch Loss: 2.306326389312744  
Epoch [2/7], Batch Loss: 2.5579192638397217  
Epoch [2/7], Average Training Loss: 2.4245523611704507

Evaluating: 100%|██████████| 19/19 [05:40<00:00, 17.94s/it]

Validation Metrics after Epoch 2: {'ROUGE-L': 0.344430507047284}

Epoch [3/7], Batch Loss: 2.1594936847686768  
Epoch [3/7], Batch Loss: 2.2254562377929688  
Epoch [3/7], Batch Loss: 2.238584041595459  
Epoch [3/7], Batch Loss: 2.3106191158294678  
Epoch [3/7], Batch Loss: 2.3417437076568604  
Epoch [3/7], Batch Loss: 2.27418851852417  
Epoch [3/7], Batch Loss: 2.168013572692871  
Epoch [3/7], Batch Loss: 2.4113423824310303

Epoch [3/7], Batch Loss: 2.413123607635498  
Epoch [3/7], Batch Loss: 2.1674985885620117  
Epoch [3/7], Batch Loss: 2.2781403064727783  
Epoch [3/7], Batch Loss: 2.450666666030884  
Epoch [3/7], Batch Loss: 2.333395004272461  
Epoch [3/7], Batch Loss: 2.5388786792755127  
Epoch [3/7], Batch Loss: 2.3976004123687744  
Epoch [3/7], Batch Loss: 2.227501392364502  
Epoch [3/7], Batch Loss: 2.392740249633789  
Epoch [3/7], Batch Loss: 2.5090410709381104  
Epoch [3/7], Batch Loss: 2.302180051803589  
Epoch [3/7], Batch Loss: 2.5127153396606445  
Epoch [3/7], Batch Loss: 2.105900526046753  
Epoch [3/7], Batch Loss: 2.37640380859375  
Epoch [3/7], Batch Loss: 2.359524965286255  
Epoch [3/7], Batch Loss: 2.455698013305664  
Epoch [3/7], Batch Loss: 2.440885305404663  
Epoch [3/7], Batch Loss: 2.180384874343872  
Epoch [3/7], Batch Loss: 2.5481936931610107  
Epoch [3/7], Batch Loss: 2.4100120067596436  
Epoch [3/7], Batch Loss: 2.3338232040405273  
Epoch [3/7], Batch Loss: 2.1490001678466797  
Epoch [3/7], Batch Loss: 2.313507556915283  
Epoch [3/7], Batch Loss: 2.412898063659668  
Epoch [3/7], Batch Loss: 2.318373680114746  
Epoch [3/7], Batch Loss: 2.3067574501037598  
Epoch [3/7], Batch Loss: 2.3636155128479004  
Epoch [3/7], Batch Loss: 2.3872010707855225  
Epoch [3/7], Batch Loss: 2.4502766132354736  
Epoch [3/7], Batch Loss: 2.4639973640441895  
Epoch [3/7], Batch Loss: 2.333822727203369  
Epoch [3/7], Batch Loss: 2.081879138946533  
Epoch [3/7], Batch Loss: 2.237070322036743  
Epoch [3/7], Batch Loss: 2.232893228530884  
Epoch [3/7], Batch Loss: 2.285958766937256  
Epoch [3/7], Batch Loss: 2.4619176387786865  
Epoch [3/7], Batch Loss: 2.2604429721832275  
Epoch [3/7], Batch Loss: 2.2655153274536133  
Epoch [3/7], Batch Loss: 2.482696771621704  
Epoch [3/7], Batch Loss: 2.1740479469299316  
Epoch [3/7], Batch Loss: 2.2022299766540527  
Epoch [3/7], Batch Loss: 2.5882089138031006  
Epoch [3/7], Batch Loss: 2.41127872467041  
Epoch [3/7], Batch Loss: 2.2348477840423584  
Epoch [3/7], Batch Loss: 2.2224011421203613  
Epoch [3/7], Batch Loss: 2.337162971496582  
Epoch [3/7], Batch Loss: 2.255171537399292  
Epoch [3/7], Batch Loss: 2.2868802547454834  
Epoch [3/7], Batch Loss: 2.2586216926574707



Epoch [3/7], Batch Loss: 2.3802683353424072  
Epoch [3/7], Batch Loss: 2.3150317668914795  
Epoch [3/7], Batch Loss: 2.4159576892852783  
Epoch [3/7], Batch Loss: 2.4201531410217285  
Epoch [3/7], Batch Loss: 2.0149495601654053  
Epoch [3/7], Batch Loss: 2.3650448322296143  
Epoch [3/7], Batch Loss: 2.2265522480010986  
Epoch [3/7], Batch Loss: 2.4751734733581543  
Epoch [3/7], Batch Loss: 2.3235912322998047  
Epoch [3/7], Batch Loss: 2.230715751647949  
Epoch [3/7], Batch Loss: 2.383399248123169  
Epoch [3/7], Batch Loss: 2.320690393447876  
Epoch [3/7], Batch Loss: 2.3636586666107178  
Epoch [3/7], Batch Loss: 2.1514906883239746  
Epoch [3/7], Batch Loss: 2.358126401901245  
Epoch [3/7], Batch Loss: 2.4013311862945557  
Epoch [3/7], Batch Loss: 2.4799017906188965  
Epoch [3/7], Batch Loss: 2.445903778076172  
Epoch [3/7], Average Training Loss: 2.32952486038208

Evaluating: 100%|██████████| 19/19 [05:39<00:00, 17.85s/it]

Validation Metrics after Epoch 3: {'ROUGE-L': 0.345049722565402}

Epoch [4/7], Batch Loss: 2.332324504852295  
Epoch [4/7], Batch Loss: 1.9613622426986694  
Epoch [4/7], Batch Loss: 2.4344804286956787  
Epoch [4/7], Batch Loss: 2.2704036235809326  
Epoch [4/7], Batch Loss: 2.047391653060913  
Epoch [4/7], Batch Loss: 2.4810235500335693  
Epoch [4/7], Batch Loss: 2.0178885459899902  
Epoch [4/7], Batch Loss: 2.5367531776428223  
Epoch [4/7], Batch Loss: 2.224294424057007  
Epoch [4/7], Batch Loss: 2.3060758113861084  
Epoch [4/7], Batch Loss: 1.9798797369003296  
Epoch [4/7], Batch Loss: 2.010374069213867  
Epoch [4/7], Batch Loss: 2.366204023361206  
Epoch [4/7], Batch Loss: 2.5163962841033936  
Epoch [4/7], Batch Loss: 2.3553290367126465  
Epoch [4/7], Batch Loss: 2.0656449794769287  
Epoch [4/7], Batch Loss: 2.4845731258392334  
Epoch [4/7], Batch Loss: 2.170170783996582  
Epoch [4/7], Batch Loss: 2.3512303829193115  
Epoch [4/7], Batch Loss: 2.300783395767212  
Epoch [4/7], Batch Loss: 2.099226951599121  
Epoch [4/7], Batch Loss: 2.3625922203063965  
Epoch [4/7], Batch Loss: 2.441337823867798  
Epoch [4/7], Batch Loss: 2.4301116466522217  
Epoch [4/7], Batch Loss: 2.3263907432556152  
Epoch [4/7], Batch Loss: 1.983954906463623  
Epoch [4/7], Batch Loss: 2.226391315460205

Epoch [4/7], Batch Loss: 2.379263162612915  
Epoch [4/7], Batch Loss: 2.207843542098999  
Epoch [4/7], Batch Loss: 2.398375988006592  
Epoch [4/7], Batch Loss: 2.414604663848877  
Epoch [4/7], Batch Loss: 2.309366226196289  
Epoch [4/7], Batch Loss: 2.2828786373138428  
Epoch [4/7], Batch Loss: 2.3672001361846924  
Epoch [4/7], Batch Loss: 2.266467571258545  
Epoch [4/7], Batch Loss: 2.1704819202423096  
Epoch [4/7], Batch Loss: 2.212552070617676  
Epoch [4/7], Batch Loss: 2.526765823364258  
Epoch [4/7], Batch Loss: 2.1123907566070557  
Epoch [4/7], Batch Loss: 2.200368642807007  
Epoch [4/7], Batch Loss: 2.364555597305298  
Epoch [4/7], Batch Loss: 2.289813756942749  
Epoch [4/7], Batch Loss: 2.570289134979248  
Epoch [4/7], Batch Loss: 2.158773899078369  
Epoch [4/7], Batch Loss: 2.2397894859313965  
Epoch [4/7], Batch Loss: 2.5088140964508057  
Epoch [4/7], Batch Loss: 1.9906984567642212  
Epoch [4/7], Batch Loss: 2.2286391258239746  
Epoch [4/7], Batch Loss: 2.2493960857391357  
Epoch [4/7], Batch Loss: 2.199307918548584  
Epoch [4/7], Batch Loss: 2.018333673477173  
Epoch [4/7], Batch Loss: 2.325873851776123  
Epoch [4/7], Batch Loss: 2.3053174018859863  
Epoch [4/7], Batch Loss: 1.8854585886001587  
Epoch [4/7], Batch Loss: 2.4385743141174316  
Epoch [4/7], Batch Loss: 2.4265918731689453  
Epoch [4/7], Batch Loss: 2.214899778366089  
Epoch [4/7], Batch Loss: 2.3062381744384766  
Epoch [4/7], Batch Loss: 1.9436981678009033  
Epoch [4/7], Batch Loss: 2.579129695892334  
Epoch [4/7], Batch Loss: 2.1832504272460938  
Epoch [4/7], Batch Loss: 2.3401811122894287  
Epoch [4/7], Batch Loss: 2.2009291648864746  
Epoch [4/7], Batch Loss: 2.140896797180176  
Epoch [4/7], Batch Loss: 2.1954007148742676  
Epoch [4/7], Batch Loss: 2.220191240310669  
Epoch [4/7], Batch Loss: 2.014578104019165  
Epoch [4/7], Batch Loss: 2.191800355911255  
Epoch [4/7], Batch Loss: 2.5481059551239014  
Epoch [4/7], Batch Loss: 1.6948726177215576  
Epoch [4/7], Batch Loss: 2.4542059898376465  
Epoch [4/7], Batch Loss: 2.120124340057373  
Epoch [4/7], Batch Loss: 2.225001335144043  
Epoch [4/7], Batch Loss: 2.4827544689178467  
Epoch [4/7], Batch Loss: 2.350822687149048  
Epoch [4/7], Average Training Loss: 2.260512758890788

Evaluating: 100%|██████████| 19/19 [05:44<00:00, 18.13s/it]

Validation Metrics after Epoch 4: {'ROUGE-L': 0.3317802685420676}

Epoch [5/7], Batch Loss: 2.312358856201172  
Epoch [5/7], Batch Loss: 2.3025755882263184  
Epoch [5/7], Batch Loss: 1.948441982269287  
Epoch [5/7], Batch Loss: 2.246487617492676  
Epoch [5/7], Batch Loss: 2.4745826721191406  
Epoch [5/7], Batch Loss: 2.2634265422821045  
Epoch [5/7], Batch Loss: 2.0459282398223877  
Epoch [5/7], Batch Loss: 2.197880983352661  
Epoch [5/7], Batch Loss: 2.2163000106811523  
Epoch [5/7], Batch Loss: 2.130763292312622  
Epoch [5/7], Batch Loss: 2.293497085571289  
Epoch [5/7], Batch Loss: 2.124683380126953  
Epoch [5/7], Batch Loss: 2.2604856491088867  
Epoch [5/7], Batch Loss: 2.1804897785186768  
Epoch [5/7], Batch Loss: 2.147610664367676  
Epoch [5/7], Batch Loss: 2.0365312099456787  
Epoch [5/7], Batch Loss: 2.1277987957000732  
Epoch [5/7], Batch Loss: 2.0480313301086426  
Epoch [5/7], Batch Loss: 1.9922146797180176  
Epoch [5/7], Batch Loss: 2.1892406940460205  
Epoch [5/7], Batch Loss: 2.080991744995117  
Epoch [5/7], Batch Loss: 2.123445987701416  
Epoch [5/7], Batch Loss: 2.134263038635254  
Epoch [5/7], Batch Loss: 2.121415138244629  
Epoch [5/7], Batch Loss: 2.3681704998016357  
Epoch [5/7], Batch Loss: 2.2607390880584717  
Epoch [5/7], Batch Loss: 2.414848804473877  
Epoch [5/7], Batch Loss: 1.9840725660324097  
Epoch [5/7], Batch Loss: 2.1422922611236572  
Epoch [5/7], Batch Loss: 2.1246800422668457  
Epoch [5/7], Batch Loss: 2.103226661682129  
Epoch [5/7], Batch Loss: 2.186436176300049  
Epoch [5/7], Batch Loss: 2.37707257270813  
Epoch [5/7], Batch Loss: 2.3411197662353516  
Epoch [5/7], Batch Loss: 2.238363027572632  
Epoch [5/7], Batch Loss: 2.4496703147888184  
Epoch [5/7], Batch Loss: 1.985742211341858  
Epoch [5/7], Batch Loss: 2.178551197052002  
Epoch [5/7], Batch Loss: 2.3275115489959717  
Epoch [5/7], Batch Loss: 2.1327927112579346  
Epoch [5/7], Batch Loss: 2.1745541095733643  
Epoch [5/7], Batch Loss: 1.800525188446045  
Epoch [5/7], Batch Loss: 2.2693941593170166  
Epoch [5/7], Batch Loss: 2.0052340030670166  
Epoch [5/7], Batch Loss: 2.260847330093384  
Epoch [5/7], Batch Loss: 2.301740884780884  
Epoch [5/7], Batch Loss: 2.212573289871216

```
Epoch [5/7], Batch Loss: 2.1069445610046387
Epoch [5/7], Batch Loss: 2.1007611751556396
Epoch [5/7], Batch Loss: 1.626859188079834
Epoch [5/7], Batch Loss: 2.037775754928589
Epoch [5/7], Batch Loss: 2.160681962966919
Epoch [5/7], Batch Loss: 2.0486199855804443
Epoch [5/7], Batch Loss: 2.2519936561584473
Epoch [5/7], Batch Loss: 2.2492048740386963
Epoch [5/7], Batch Loss: 2.3523669242858887
Epoch [5/7], Batch Loss: 2.0610201358795166
Epoch [5/7], Batch Loss: 2.2708969116210938
Epoch [5/7], Batch Loss: 2.157914638519287
Epoch [5/7], Batch Loss: 2.434263229370117
Epoch [5/7], Batch Loss: 2.2026724815368652
Epoch [5/7], Batch Loss: 2.544494867324829
Epoch [5/7], Batch Loss: 2.1853957176208496
Epoch [5/7], Batch Loss: 2.1029481887817383
Epoch [5/7], Batch Loss: 2.3391783237457275
Epoch [5/7], Batch Loss: 2.367377281188965
Epoch [5/7], Batch Loss: 2.312711238861084
Epoch [5/7], Batch Loss: 2.532625675201416
Epoch [5/7], Batch Loss: 2.1843369007110596
Epoch [5/7], Batch Loss: 2.0345985889434814
Epoch [5/7], Batch Loss: 2.6099116802215576
Epoch [5/7], Batch Loss: 2.1024951934814453
Epoch [5/7], Batch Loss: 2.210735559463501
Epoch [5/7], Batch Loss: 1.9930121898651123
Epoch [5/7], Batch Loss: 2.1938745975494385
Epoch [5/7], Average Training Loss: 2.192203664779663
```

```
Evaluating: 100%|██████████| 19/19 [05:38<00:00, 17.80s/it]
```

Validation Metrics after Epoch 5: {'ROUGE-L': 0.33798108652954467}

Epoch [6/7], Batch Loss: 2.0221457481384277

Epoch [6/7], Batch Loss: 1.918637752532959

Epoch [6/7], Batch Loss: 2.465228319168091

[illegible]

```
<ipython-input-19-200128bbeb7c> in <cell line: 35>()
```

```

62
63         # Set max_length explicitly for BLIP
generation to avoid inconsistent lengths
--> 64         outputs = blip_model.generate(*inputs,
max_length=200) # Limit caption length to 150 tokens
65         caption = blip_processor.decode(outputs[0],
skip_special_tokens=True)
66         single_captions.append(caption)

```

```

/usr/local/lib/python3.10/dist-packages/torch/utils/_contextlib.py in
decorate_context(*args, **kwargs)
    114     def decorate_context(*args, **kwargs):
    115         with ctx_factory():
--> 116             return func(*args, **kwargs)
    117
    118     return decorate_context

```

```

/usr/local/lib/python3.10/dist-packages/transformers/models/blip/model
ing_blip.py in generate(self, pixel_values, input_ids, attention_mask,
interpolate_pos_encoding, **generate_kwargs)
    1185         attention_mask = attention_mask[:, :-1] if
attention_mask is not None else None
    1186
-> 1187         outputs = self.text_decoder.generate(
    1188             input_ids=input_ids[:, :-1],
    1189             eos_token_id=self.config.text_config.sep_token_id,

```

```

/usr/local/lib/python3.10/dist-packages/torch/utils/_contextlib.py in
decorate_context(*args, **kwargs)
    114     def decorate_context(*args, **kwargs):
    115         with ctx_factory():
--> 116             return func(*args, **kwargs)
    117
    118     return decorate_context

```

```

/usr/local/lib/python3.10/dist-packages/transformers/generation/utils.
py in generate(self, inputs, generation_config, logits_processor,
stopping_criteria, prefix_allowed_tokens_fn, synced_gpus,
assistant_model, streamer, negative_prompt_ids,
negative_prompt_attention_mask, **kwargs)
    2022
    2023         # 13. run sample (it degenerates to greedy search
when `generation_config.do_sample=False`)
-> 2024         result = self._sample(
    2025             input_ids,
    2026             logits_processor=prepared_logits_processor,

```

```

/usr/local/lib/python3.10/dist-packages/transformers/generation/utils.
py in _sample(self, input_ids, logits_processor, stopping_criteria,
generation_config, synced_gpus, streamer, logits_warper,
**model_kwargs)
    2980
    2981         # forward pass to get next token
-> 2982         outputs = self(**model_inputs, return_dict=True)
    2983
    2984         if synced_gpus and this_peer_finished:

```

```

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in

```

```

_wrapped_call_impl(self, *args, **kwargs)
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735         else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None

/usr/local/lib/python3.10/dist-packages/transformers/models/blip/model
ing_blip_text.py in forward(self, input_ids, attention_mask,
position_ids, head_mask, inputs_embeds, encoder_hidden_states,
encoder_attention_mask, labels, past_key_values, use_cache,
output_attentions, output_hidden_states, return_dict, return_logits,
is_decoder, reduction)
869         use_cache = False
870
--> 871         outputs = self.bert(
872             input_ids,
873             attention_mask=attention_mask,

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_wrapped_call_impl(self, *args, **kwargs)
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735         else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None

```

```
/usr/local/lib/python3.10/dist-packages/transformers/models/blip/modeling_blip_text.py in forward(self, input_ids, attention_mask, position_ids, head_mask, inputs_embeds, encoder_embeds, encoder_hidden_states, encoder_attention_mask, past_key_values, use_cache, output_attentions, output_hidden_states, return_dict, is_decoder)
```

```
780         embedding_output = encoder_embeds
781
--> 782         encoder_outputs = self.encoder(
783             embedding_output,
784             attention_mask=extended_attention_mask,
```

```
/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in _wrapped_call_impl(self, *args, **kwargs)
```

```
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735     else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code
```

```
/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in _call_impl(self, *args, **kwargs)
```

```
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None
```

```
/usr/local/lib/python3.10/dist-packages/transformers/models/blip/modeling_blip_text.py in forward(self, hidden_states, attention_mask, head_mask, encoder_hidden_states, encoder_attention_mask, past_key_values, use_cache, output_attentions, output_hidden_states, return_dict)
```

```
434         )
435     else:
--> 436         layer_outputs = layer_module(
437             hidden_states,
438             attention_mask,
```

```
/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in _wrapped_call_impl(self, *args, **kwargs)
```

```
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735     else:
-> 1736         return self._call_impl(*args, **kwargs)
```

```

1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None

/usr/local/lib/python3.10/dist-packages/transformers/models/blip/model
ing_blip_text.py in forward(self, hidden_states, attention_mask,
head_mask, encoder_hidden_states, encoder_attention_mask,
past_key_value, output_attentions)
356
357         if encoder_hidden_states is not None:
--> 358             cross_attention_outputs = self.crossattention(
359                 attention_output,
360                 attention_mask,

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_wrapped_call_impl(self, *args, **kwargs)
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735     else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None

/usr/local/lib/python3.10/dist-packages/transformers/models/blip/model
ing_blip_text.py in forward(self, hidden_states, attention_mask,
head_mask, encoder_hidden_states, encoder_attention_mask,
past_key_value, output_attentions)
273         output_attentions: Optional[bool] = False,
274     ) -> Tuple[torch.Tensor]:

```



```

--> 275         self_outputs = self.self(
276             hidden_states,
277             attention_mask,

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_wrapped_call_impl(self, *args, **kwargs)
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735     else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks
1746         or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747         return forward_call(*args, **kwargs)
1748
1749         result = None

/usr/local/lib/python3.10/dist-packages/transformers/models/blip/model
ing_blip_text.py in forward(self, hidden_states, attention_mask,
head_mask, encoder_hidden_states, encoder_attention_mask,
past_key_value, output_attentions)
160         if is_cross_attention:
161             key_layer =
self.transpose_for_scores(self.key(encoder_hidden_states))
--> 162             value_layer =
self.transpose_for_scores(self.value(encoder_hidden_states))
163             attention_mask = encoder_attention_mask
164             elif past_key_value is not None:

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_wrapped_call_impl(self, *args, **kwargs)
1734         return self._compiled_call_impl(*args, **kwargs)
# type: ignore[misc]
1735     else:
-> 1736         return self._call_impl(*args, **kwargs)
1737
1738     # torchrec tests the code consistency with the following
code

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_call_impl(self, *args, **kwargs)
1745         or _global_backward_pre_hooks or
_global_backward_hooks

```

```

1746             or _global_forward_hooks or
_global_forward_pre_hooks):
-> 1747             return forward_call(*args, **kwargs)
1748
1749             result = None

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/linear.py in
forward(self, input)
123
124     def forward(self, input: Tensor) -> Tensor:
--> 125         return F.linear(input, self.weight, self.bias)
126
127     def extra_repr(self) -> str:

```

KeyboardInterrupt:

```

rougeL_scores = [0.3243333918983624, 0.344430507047284,
0.345049722565402, 0.3317802685420676, 0.33798108652954467]
training_losses = [2.621926883061727, 2.4245523611704507,
2.32952486038208, 2.260512758890788, 2.192203664779663 ]

```

```
import torch
```

```
# Define local save paths
```

```
blip_save_path = 'blip_model.pth'
```

```
t5_save_path = 't5_model.pth'
```

```
# Save the state_dict of each model
```

```
torch.save(blip_model.state_dict(), blip_save_path)
```

```
torch.save(t5_model.state_dict(), t5_save_path)
```

```
print("Models have been saved locally.")
```

Models have been saved locally.

```
# Evaluate after each epoch
```

```
metrics = evaluate_model_training(blip_model, t5_model,
```

```
blip_processor, t5_tokenizer, train_dataloader, device)
```

```
# Print evaluation metrics for each epoch
```

```
print(f"Training Metric after Epoch {epoch+1}: {metrics}")
```

```
Evaluating: 93%|██████████| 70/75 [21:35<01:34, 18.85s/it]
```

```
# After training each epoch, evaluate the model on the validation set
```

```
metrics = evaluate_model(blip_model, t5_model, blip_processor,
```

```
t5_tokenizer, val_dataloader, device)
```

```
print(f"Validation Metrics after Epoch {epoch+1}: {metrics}")
```

```
Evaluating: 100%|██████████| 19/19 [05:36<00:00, 17.72s/it]
```

```
{"model_id": "f75d3086443c4945ad9ec5528628aee9", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f0da07ee1eb74b68bcd1d52ec89386d3d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d0d8485d09bb475981e2631b4c757bf0", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "ae383ae5940d4482b3a3426e9c580d5d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "625558015b684bf4a6423bb03fe819a4", "version_major": 2, "version_minor": 0}
```

```
/usr/local/lib/python3.10/dist-packages/transformers/
tokenization_utils_base.py:1601: FutureWarning:
`clean_up_tokenization_spaces` was not set. It will be set to `True`
by default. This behavior will be deprecated in transformers v4.45, and
will be then set to `False` by default. For more details check this
issue: https://github.com/huggingface/transformers/issues/31884
  warnings.warn(
```

```
{"model_id": "3c8fb8f034ad4bd5ab246b33f215ffaa", "version_major": 2, "version_minor": 0}
```

Some weights of RobertaModel were not initialized from the model checkpoint at roberta-large and are newly initialized:  
['roberta.pooler.dense.bias', 'roberta.pooler.dense.weight']  
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
calculating scores...
computing bert embedding.
```

```
{"model_id": "97634b1077a243b78141b9454cca0ecb", "version_major": 2, "version_minor": 0}
```

```
computing greedy matching.
```

```
{"model_id": "6d294038200c4e2baf391a53851c104d", "version_major": 2, "version_minor": 0}
```

```
done in 1.98 seconds, 151.82 sentences/sec
Validation Metrics after Epoch 6: {'BLEU': 0.06442781788659759,
'ROUGE-1': 0.3742467614881393, 'ROUGE-2': 0.1251457245015394, 'ROUGE-L': 0.3411455272680421, 'BERT_F1': 0.9064300656318665}
```

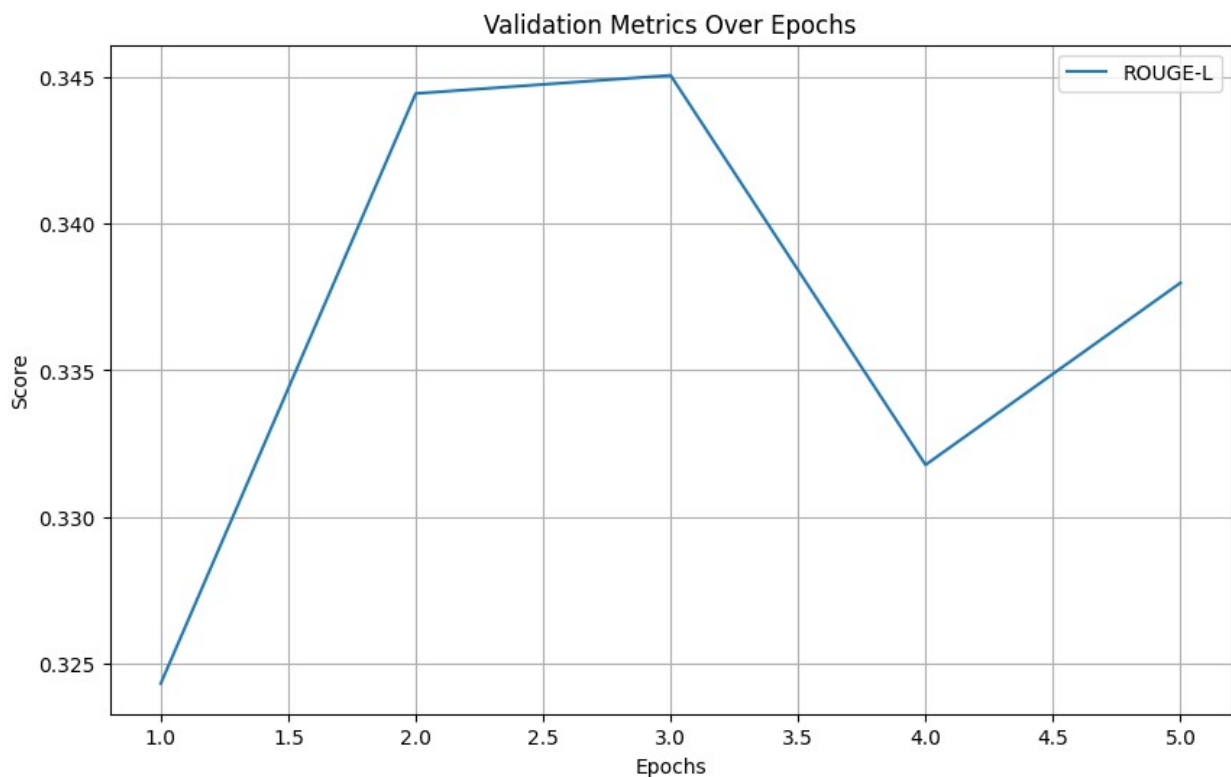
```
import matplotlib.pyplot as plt
```

```
# Plot training loss over epochs
```

```

plt.figure(figsize=(10, 6))
#plt.plot(range(1, epochs + 1), bleu_scores, label='BLEU')
#plt.plot(range(1, epochs + 1), rouge1_scores, label='ROUGE-1')
#plt.plot(range(1, epochs + 1), rouge2_scores, label='ROUGE-2')
plt.plot(range(1, 6), rougeL_scores, label='ROUGE-L')
#plt.plot(range(1, epochs + 1), meteor_scores, label='METEOR')
#plt.plot(range(1, epochs + 1), bert_f1_scores, label='BERT F1')
plt.xlabel('Epochs')
plt.ylabel('Score')
plt.title('Validation Metrics Over Epochs')
plt.legend()
plt.grid()
plt.show()

```



```


import matplotlib.pyplot as plt

# Plot training loss over epochs

plt.figure(figsize=(10, 6))
plt.plot(range(1, 6), training_losses, label='ROUGE-L')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training Loss Over Epochs')
plt.legend()

```

```
plt.grid()
plt.show()
```



```
import requests
import imageio
from PIL import Image
from io import BytesIO
import torch
from nltk.translate.bleu_score import corpus_bleu
from rouge_score import rouge_scorer
from bert_score import score as bert_score_fn
import torch.nn.functional as F

def process_and_evaluate_gif(
    gif_url,
    actual_description,
    blip_processor,
    blip_model,
    t5_tokenizer,
    t5_model,
    device='cuda' if torch.cuda.is_available() else 'cpu',
    num_frames=5,
    frame_size=(256, 256),
    max_length=128,
```

```

max_target_length=150,
num_beams=4
):
    """
    Processes an unseen GIF to generate a description and evaluates it
    against the actual description.

    Args:
        gif_url (str): URL of the GIF to process.
        actual_description (str): The ground truth description of the
        GIF.
        blip_processor: BLIP processor instance.
        blip_model: Trained BLIP model instance.
        t5_tokenizer: T5 tokenizer instance.
        t5_model: Trained T5 model instance.
        device (str): Device to run the models on ('cuda' or 'cpu').
        num_frames (int): Number of frames to extract from the GIF.
        frame_size (tuple): Desired frame size (width, height).
        max_length (int): Maximum token length for T5 input.
        max_target_length (int): Maximum token length for T5 output.
        num_beams (int): Number of beams for beam search in T5
        generation.

    Returns:
        dict: Evaluation metrics including BLEU, ROUGE-1, ROUGE-2,
        ROUGE-L, and BERT_F1.
    """
    # Step 1: Download the GIF
    try:
        response = requests.get(gif_url, timeout=10)
        response.raise_for_status()
        gif_bytes = BytesIO(response.content)
        print(f"Successfully downloaded GIF from {gif_url}")
    except requests.exceptions.RequestException as e:
        print(f"Failed to download GIF from {gif_url}: {e}")
        return None

    # Step 2: Extract frames from the GIF
    try:
        gif = imageio.mimread(gif_bytes, memtest=False)
        total_frames = len(gif)
        if total_frames == 0:
            print("No frames found in the GIF.")
            return None
        interval = max(total_frames // num_frames, 1)
        selected_frames = [gif[i] for i in range(0, total_frames,
interval)][[:num_frames]]
        print(f"Extracted {len(selected_frames)} frames from the
        GIF.")
    except Exception as e:

```

```

        print(f"Error extracting frames from GIF: {e}")
        return None

    # Step 3: Ensure each GIF has exactly num_frames by padding if
    necessary
    while len(selected_frames) < num_frames:
        if selected_frames:
            selected_frames.append(selected_frames[-1])
        else:
            # Append a blank frame if no frames are extracted
            selected_frames.append(np.zeros((frame_size[1],
frame_size[0], 3), dtype=np.uint8))
        selected_frames = selected_frames[:num_frames] # Ensure no more
        than num_frames
    print(f"Total frames after padding: {len(selected_frames)}")

    # Step 4: Convert frames to PIL Images, resize, and generate
    captions using BLIP
    captions = []
    for idx, frame in enumerate(selected_frames):
        try:
            img = Image.fromarray(frame).convert('RGB')
            img_resized = img.resize(frame_size)
            inputs = blip_processor(img_resized,
return_tensors="pt").to(device)
            with torch.no_grad():
                outputs = blip_model.generate(**inputs, max_length=50)
                caption = blip_processor.decode(outputs[0],
skip_special_tokens=True)
            captions.append(caption)
            print(f"Frame {idx+1}: {caption}")
        except Exception as e:
            print(f"Failed to generate caption for frame {idx+1}:
{e}")
        captions.append("") # Append empty string for failed
        captions

    # Step 5: Concatenate captions to form T5 input
    concatenated_captions = " ".join(captions)
    print(f"Concatenated Captions: {concatenated_captions}")

    # Step 6: Tokenize the concatenated captions for T5
    try:
        encoding = t5_tokenizer.encode_plus(
            concatenated_captions,
            add_special_tokens=True,
            max_length=max_length,
            padding='max_length',
            truncation=True,
            return_attention_mask=True,

```

```

        return_tensors='pt',
    )
    input_ids = encoding['input_ids'].to(device)
    attention_mask = encoding['attention_mask'].to(device)
except Exception as e:
    print(f"Error during tokenization: {e}")
    return None

# Step 7: Generate description using T5
try:
    t5_model.eval()
    with torch.no_grad():
        outputs = t5_model.generate(
            input_ids=input_ids,
            attention_mask=attention_mask,
            max_length=max_target_length,
            num_beams=num_beams,
            early_stopping=True
        )
    generated_description = t5_tokenizer.decode(outputs[0],
skip_special_tokens=True)
    print(f"Generated Description: {generated_description}")
except Exception as e:
    print(f"Error during T5 generation: {e}")
    return None

# Step 8: Evaluate the generated description against the actual
description
scorer = rouge_scorer.RougeScorer(['rouge1', 'rouge2', 'rougeL'],
use_stemmer=True)

# Prepare tokenized inputs
ref_tokens = actual_description.split()
hyp_tokens = generated_description.split()

# Calculate BLEU
bleu = corpus_bleu([[ref_tokens]], [hyp_tokens])

# Calculate ROUGE
scores = scorer.score(actual_description, generated_description)
rouge1 = scores['rouge1'].fmeasure
rouge2 = scores['rouge2'].fmeasure
rougeL = scores['rougeL'].fmeasure

# Calculate BERT Score
P, R, F1 = bert_score_fn(
    [' '.join(hyp_tokens)],
    [' '.join(ref_tokens)],
    lang='en',
    verbose=False

```



```

)
bert_f1 = F1.mean().item()

# Compile metrics
metrics = {
    'BLEU': bleu,
    'ROUGE-1': rouge1,
    'ROUGE-2': rouge2,
    'ROUGE-L': rougeL,
    'BERT_F1': bert_f1
}

# Display Metrics
print("\nEvaluation Metrics:")
for metric, value in metrics.items():
    print(f"{metric}: {value:.4f}")

return {
    'generated_description': generated_description,
    'actual_description': actual_description,
    'metrics': metrics
}

# Example unseen GIF
unseen_gif = {
    'url':
'https://38.media.tumblr.com/f754d72da3c6a58211c760d39dff5be3/tumblr_n
8vbpHDLh1qdzzbkol_250.gif',
    'actual_description': 'a man in a tuxedo stares as smoke rises
next to him'
}

# Process and evaluate the unseen GIF
result = process_and_evaluate_gif(
    gif_url=unseen_gif['url'],
    actual_description=unseen_gif['actual_description'],
    blip_processor=blip_processor,
    blip_model=blip_model,
    t5_tokenizer=t5_tokenizer,
    t5_model=t5_model,
    device=device,
    num_frames=5,
    frame_size=(256, 256),
    max_length=128,
    max_target_length=150,
    num_beams=4
)

# Check the result
if result:

```

```

print("\nFinal Results:")
print(f"Generated Description: {result['generated_description']}")
print(f"Actual Description: {result['actual_description']}")
print("Evaluation Metrics:")
for metric, score in result['metrics'].items():
    print(f"    {metric}: {score:.4f}")
else:
    print("Processing and evaluation failed.")

```

Successfully downloaded GIF from  
[https://38.media.tumblr.com/f754d72da3c6a58211c760d39dff5be3/tumblr\\_n8vbphDLEhlqdzzbko1\\_250.gif](https://38.media.tumblr.com/f754d72da3c6a58211c760d39dff5be3/tumblr_n8vbphDLEhlqdzzbko1_250.gif)

Extracted 5 frames from the GIF.

Total frames after padding: 5

Frame 1: a man in a tuxed suit and bow tie

Frame 2: a man in a tuxed suit and bow tie

Frame 3: a man in a tuxed suit and bow tie

Frame 4: a man in a tuxed suit and bow tie

Frame 5: a man in a tuxed suit and bow tie

Concatenated Captions: a man in a tuxed suit and bow tie a man in a  
tuxed suit and bow tie a man in a tuxed suit and bow tie a man in a  
tuxed suit and bow tie a man in a tuxed suit and bow tie

Generated Description: a man in a tuxedo wears a tuxedo with a hooded  
shirt.

Some weights of RobertaModel were not initialized from the model  
checkpoint at roberta-large and are newly initialized:

['roberta.pooler.dense.bias', 'roberta.pooler.dense.weight']

You should probably TRAIN this model on a down-stream task to be able  
to use it for predictions and inference.

Evaluation Metrics:

BLEU: 0.3170

ROUGE-1: 0.4167

ROUGE-2: 0.3636

ROUGE-L: 0.4167

BERT\_F1: 0.9107

Final Results:

Generated Description: a man in a tuxedo wears a tuxedo with a hooded  
shirt.

Actual Description: a man in a tuxedo stares as smoke rises next to  
him

Evaluation Metrics:

BLEU: 0.3170

ROUGE-1: 0.4167

ROUGE-2: 0.3636

ROUGE-L: 0.4167

BERT\_F1: 0.9107